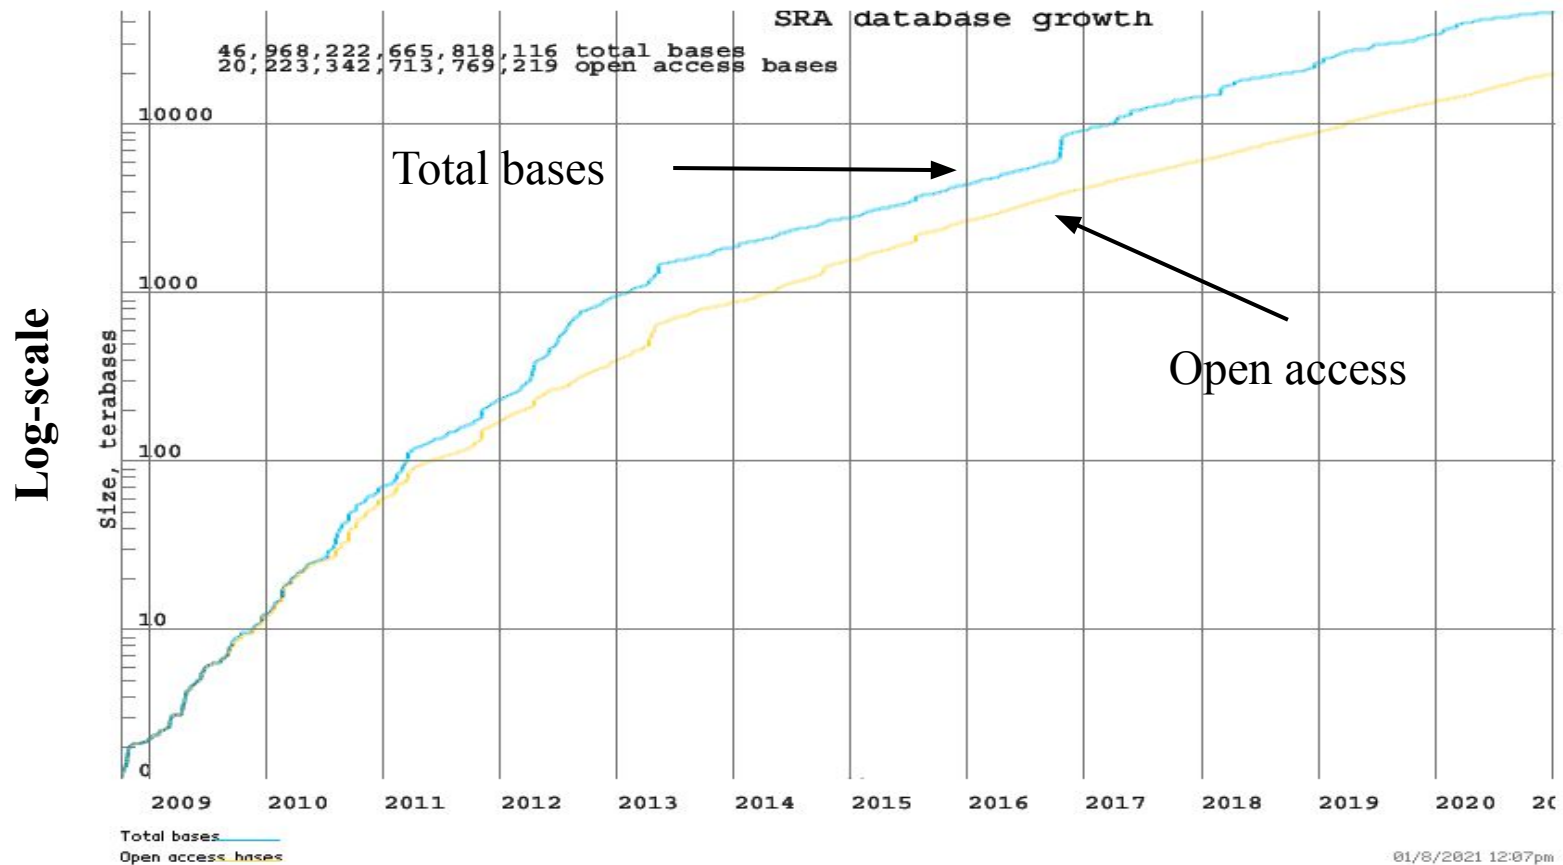


Data Science at Scale:

Scaling Up by Scaling Down and Out (to Disk)

Prashant Pandey
ppandey@berkeley.edu
Berkeley Lab/UC Berkeley

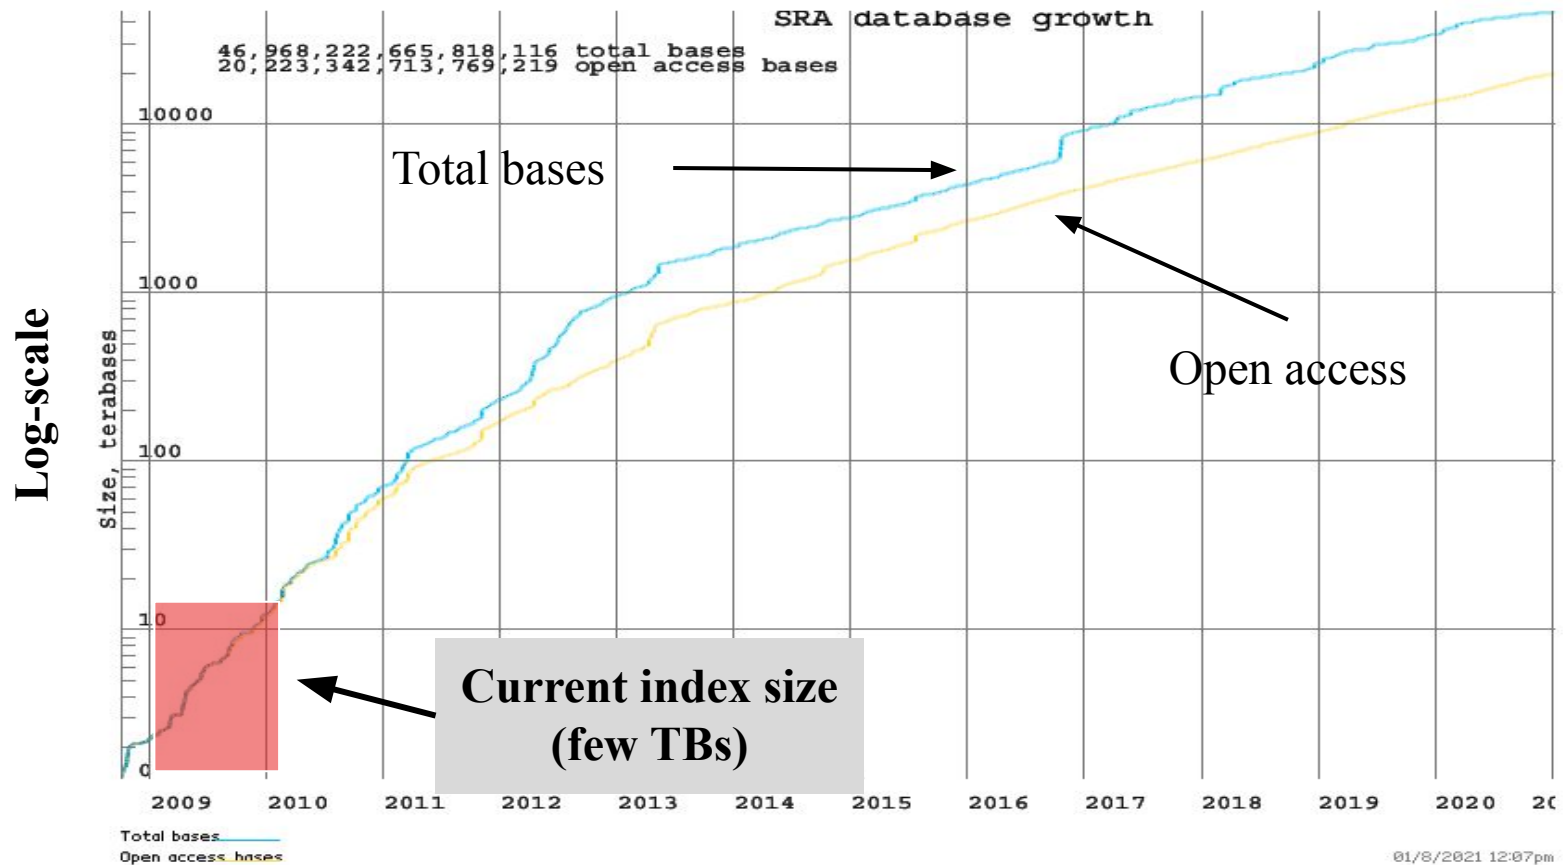
Sequence Read Archive (SRA) database growth



SRA contains a lot of *diversity information*

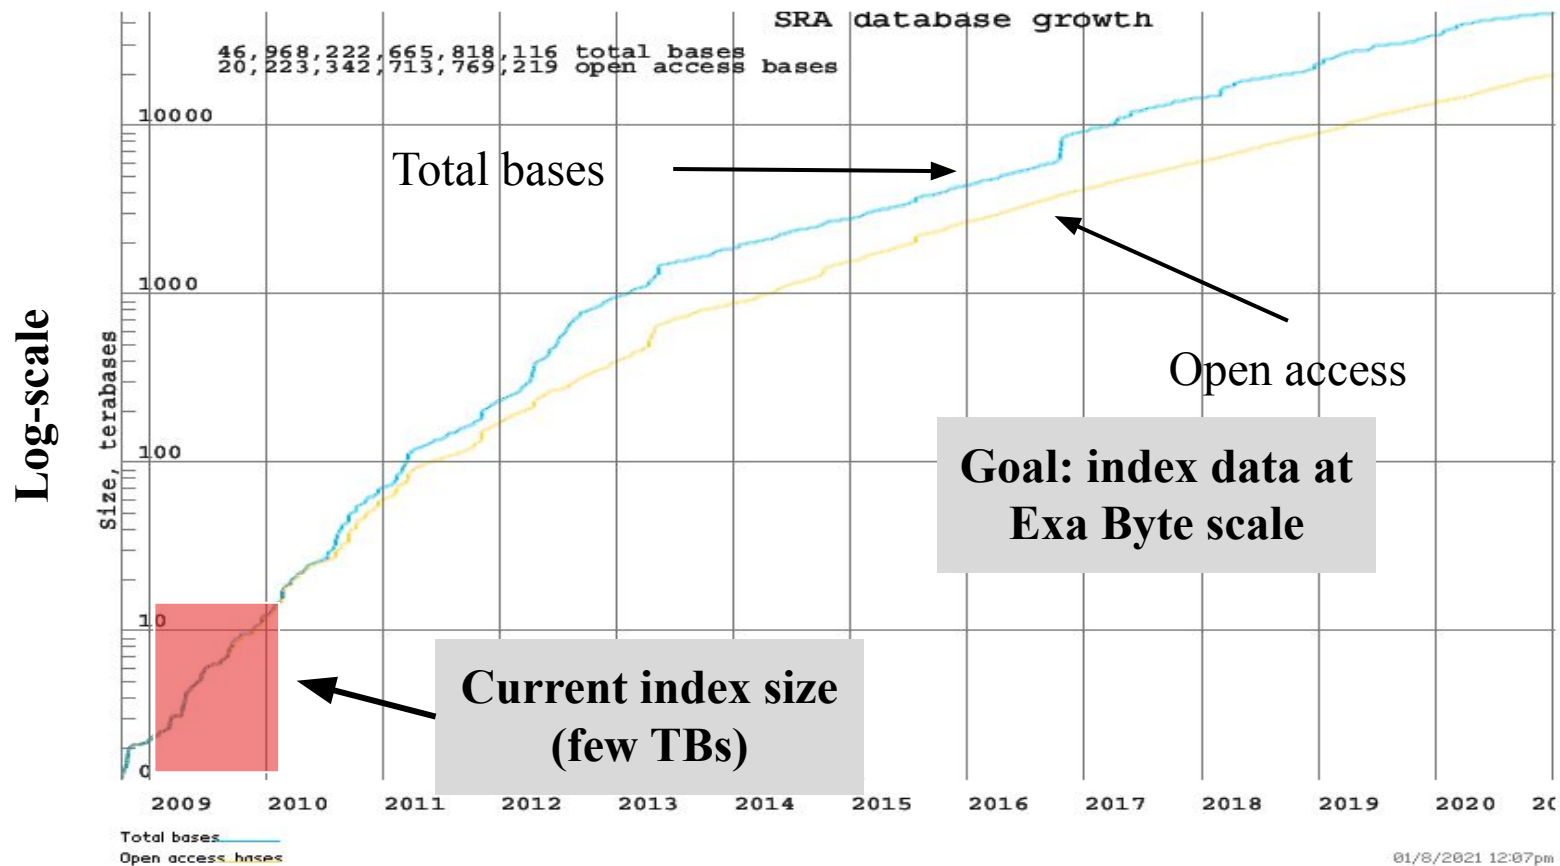
Goal: perform *sequence searches* on the database

Scalability is the bottleneck for data science



Data science applications only looking at a *small portion* of data

Scalable data systems → Scalable data science



My goal as a researcher is to build *scalable data systems* to *accelerate* and *scale data science* applications

Three approaches to handle massive data

Three approaches to handle massive data

Shrink it

Goal: make data smaller to fit in RAM

Techniques:

- Compact & succinct data structures
- Filters, e.g., Bloom, quotient, etc.

Three approaches to handle massive data

Shrink it

Goal: make data smaller to fit in RAM

Techniques:

- Compact & succinct data structures
- Filters, e.g., Bloom, quotient, etc.

Organize it

Goal: organize data in a disk-friendly way

Techniques:

- B-tree
- B⁺-tree
- LSM-tree

Three approaches to handle massive data

Shrink it

Goal: make data smaller to fit in RAM

Techniques:

- Compact & succinct data structures
- Filters, e.g., Bloom, quotient, etc.

Organize it

Goal: organize data in a disk-friendly way

Techniques:

- B-tree
- B^ε-tree
- LSM-tree

Distribute it

Goal: partition and distribute data on multiple nodes

Techniques:

- Distributed hash table
- Distributed key-value store

Research output

Data structures & Algorithms

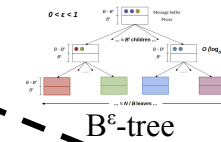
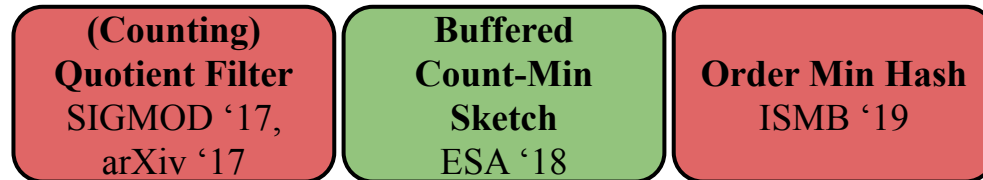
**(Counting)
Quotient Filter**
SIGMOD '17,
arXiv '17

**Buffered
Count-Min
Sketch**
ESA '18

Order Min Hash
ISMB '19

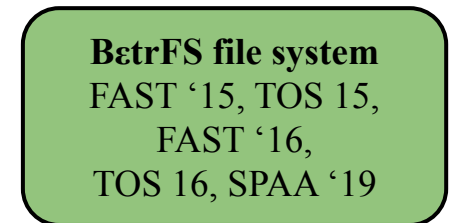
Research output

Data structures & Algorithms



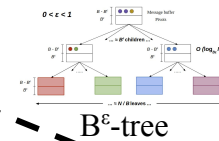
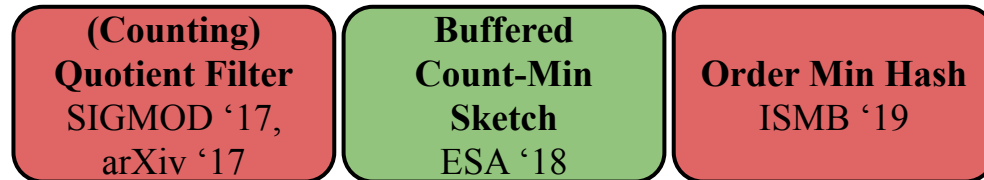
B^c-tree

File systems



Research output

Data structures & Algorithms



Computational biology

Squeakr, deBGR, Mantis, Rainbowfish, MST-Mantis
ISMB '17, WABI '17,
BIOINFORMATICS '17,
RECOMB '18, Cell Systems
'18, RECOMB '19,
JCB '20

LSM-Mantis, VaraintStore
bioRxiv '20, bioRxiv '21

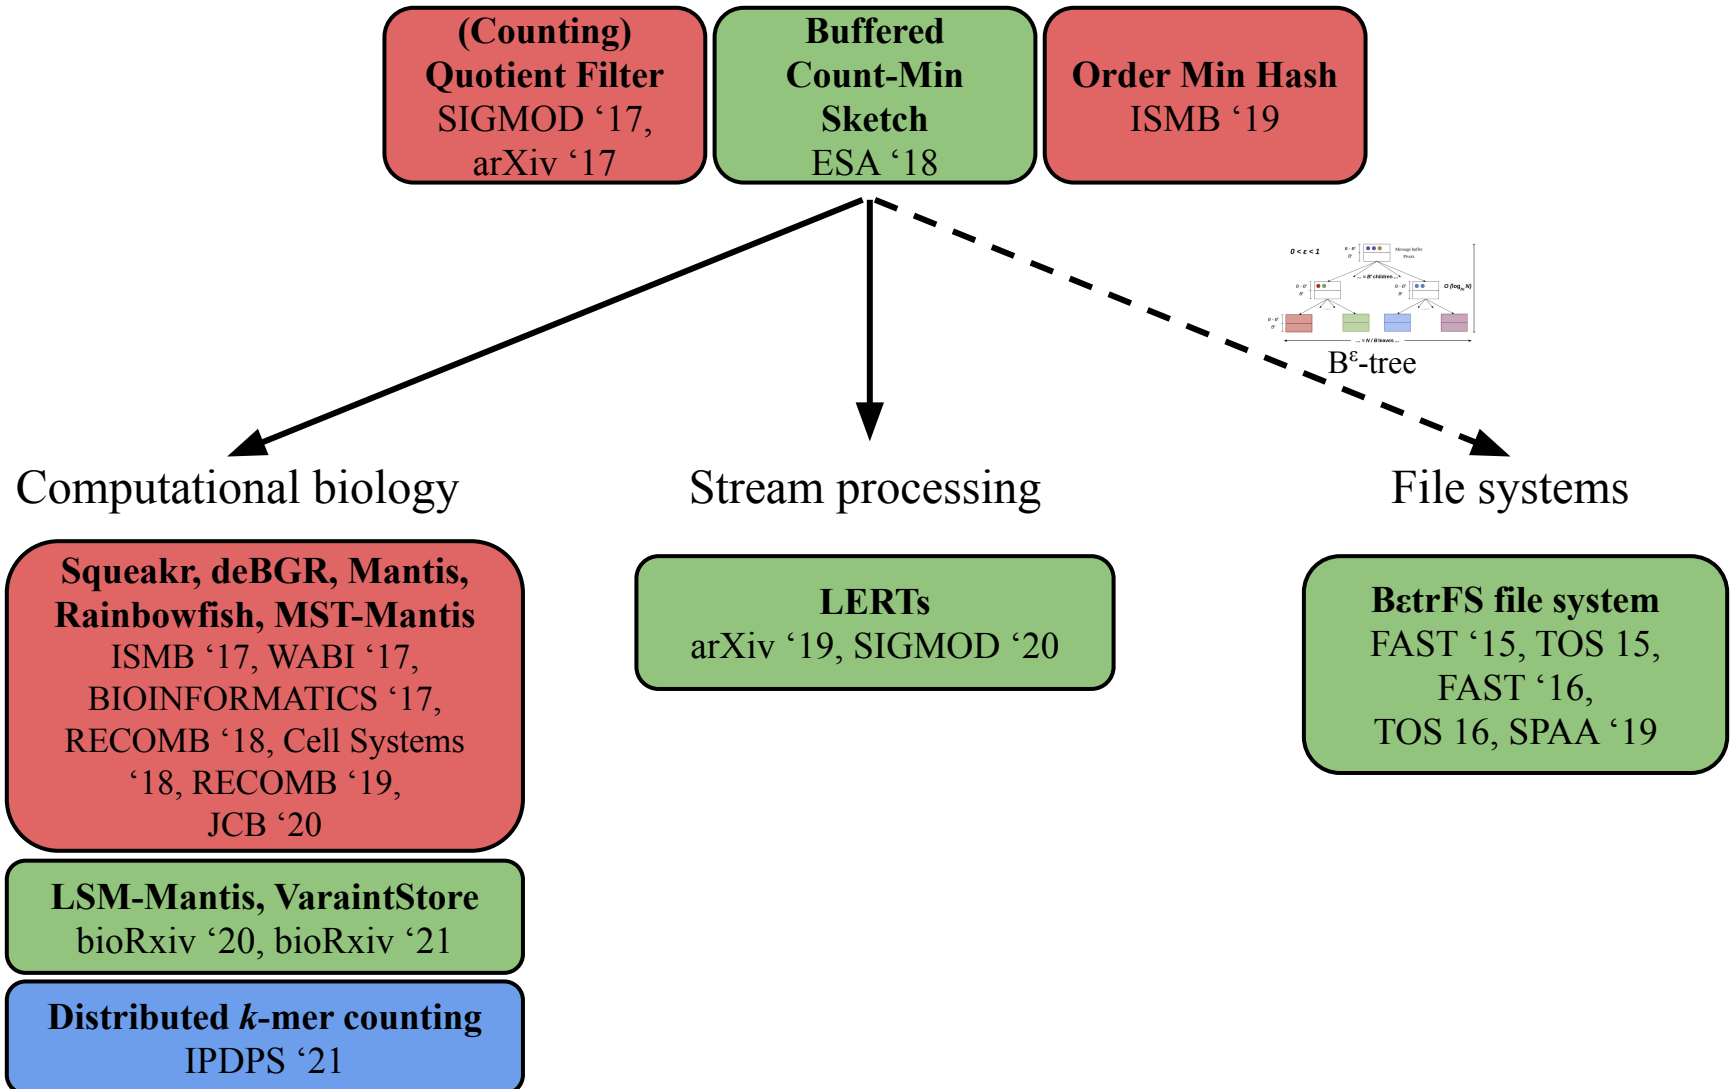
Distributed k -mer counting
IPDPS '21

File systems

B&trFS file system
FAST '15, TOS 15,
FAST '16,
TOS 16, SPAA '19

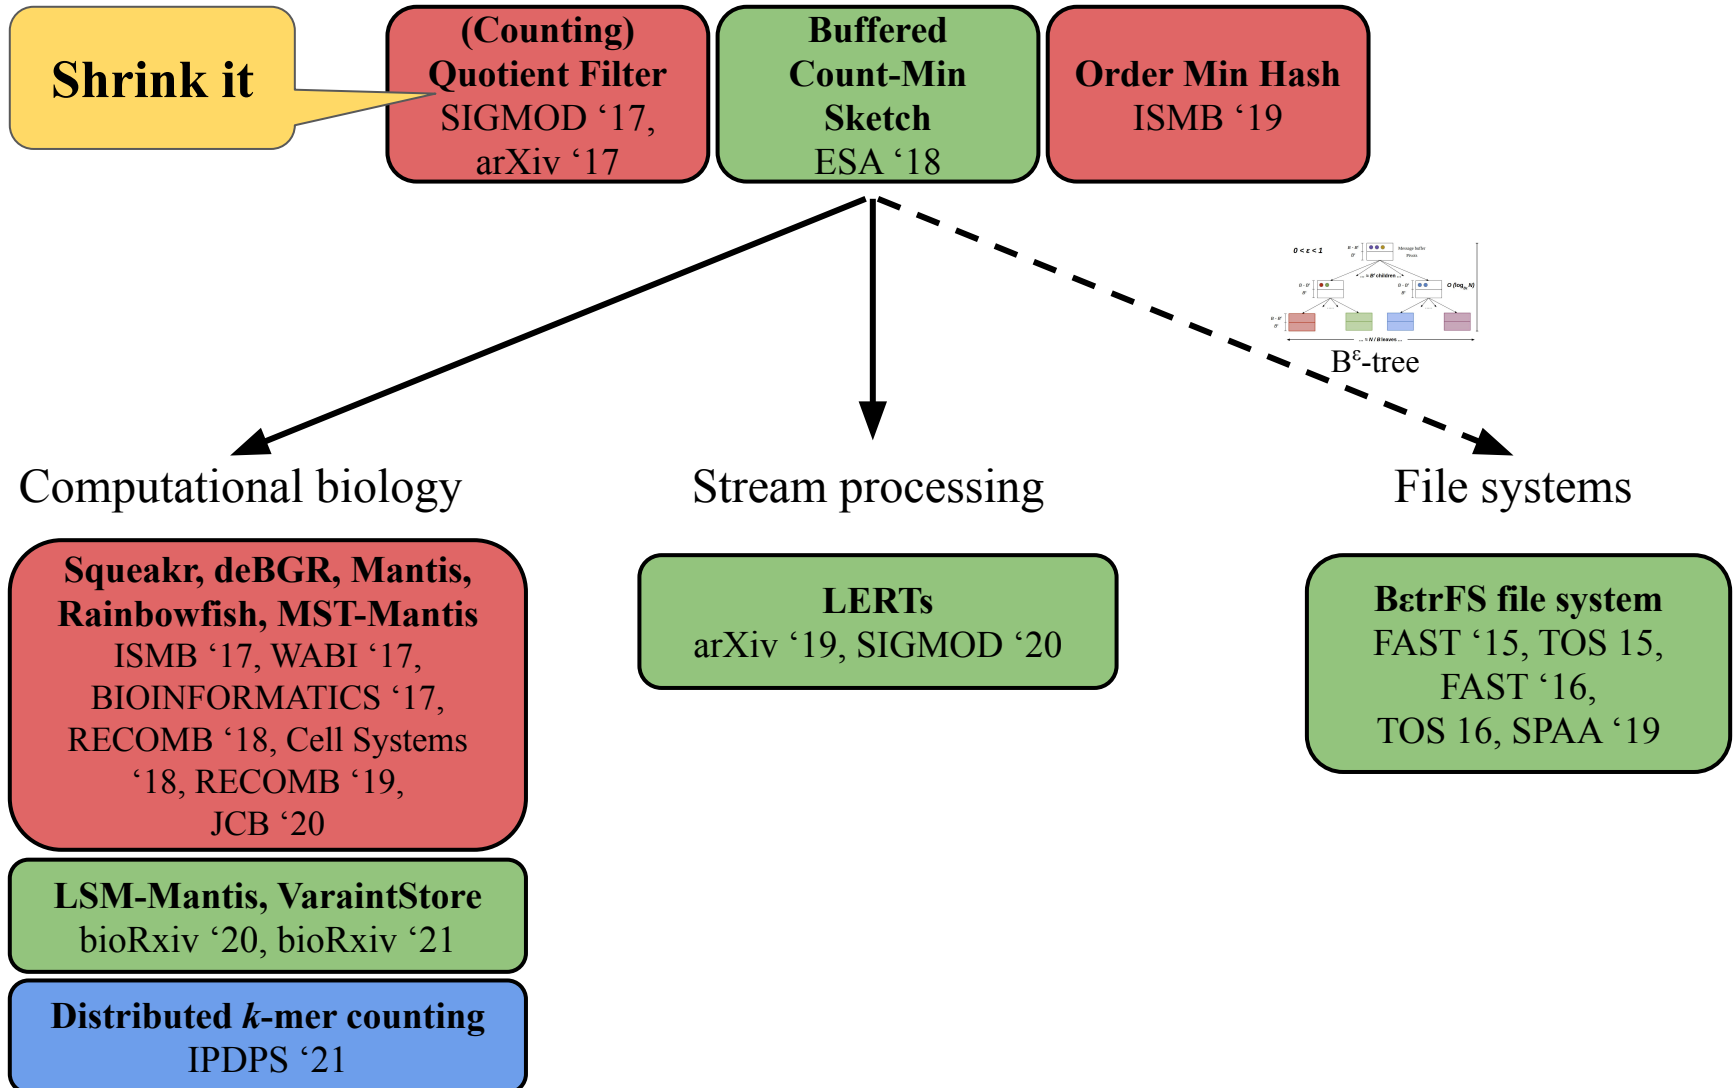
Research output

Data structures & Algorithms



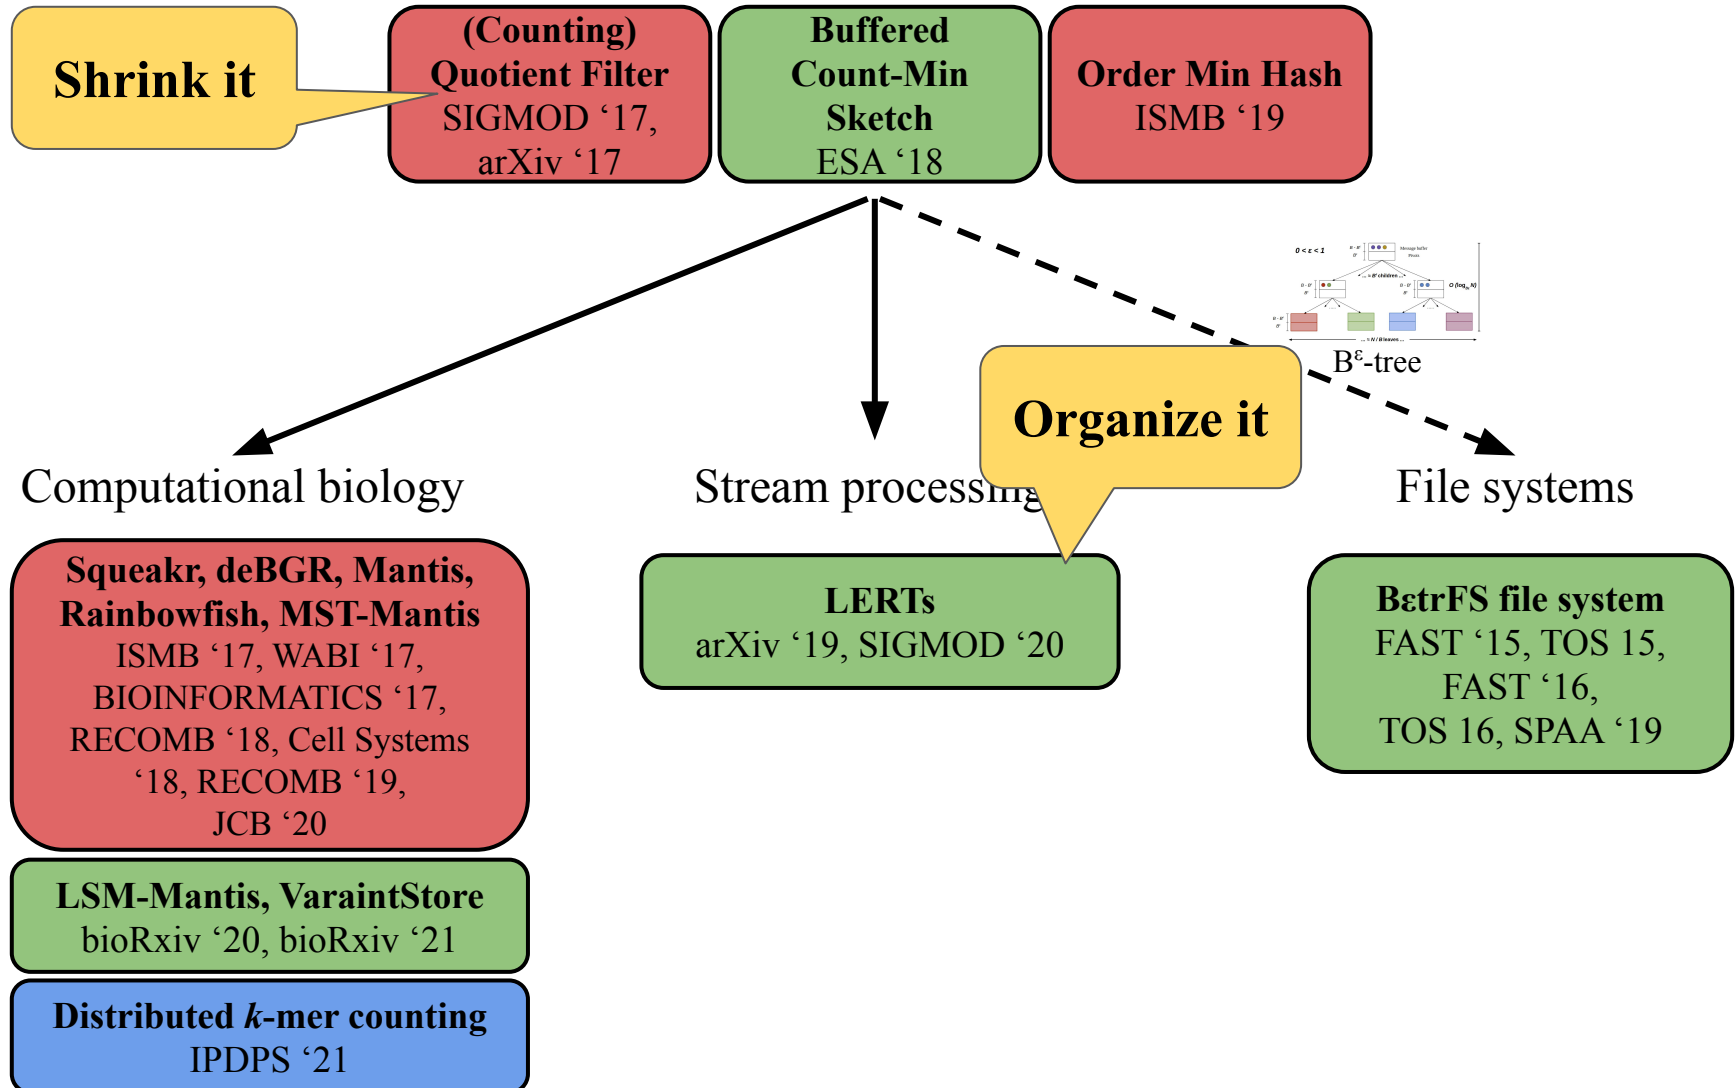
In this talk

Data structures & Algorithms




In this talk

Data structures & Algorithms



Dictionary data structure

A dictionary maintains a set S from universe U .



membership(a): ✓

membership(b): ✗


membership(c): ✓

membership(d): ✗

A dictionary supports membership queries on S .

Filter data structure

A filter is an *approximate* dictionary.



membership(a): ✓


membership(b): ✗

membership(c): ✓

membership(d): ✓ 🙅 **false positive**

A filter supports approximate membership queries on S .

A filter guarantees a false-positive rate ε

if $q \in S$, return  with probability 1 **true positive**


if $q \notin S$, return $\left\{ \begin{array}{ll} \text{✗} & \text{with probability } \leq 1 - \varepsilon \quad \text{true negative} \\ \text{✓} & \text{with probability } \leq \varepsilon \quad \text{false positive} \end{array} \right.$

one-sided
errors


False-positive rate enables filters to be compact

$$\text{space} \geq n \log(1/\epsilon)$$

$$\text{space} = \Omega(n \log |U|)$$



Filter




Dictionary

False-positive rate enables filters to be compact

$$\text{space} \geq n \log(1/\epsilon)$$

Small



Filter

$$\text{space} = \Omega(n \log |U|)$$

Large



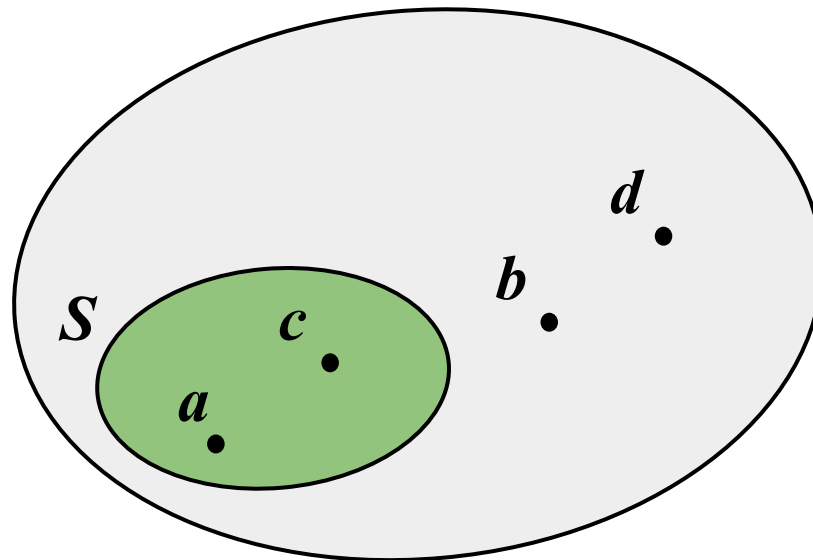
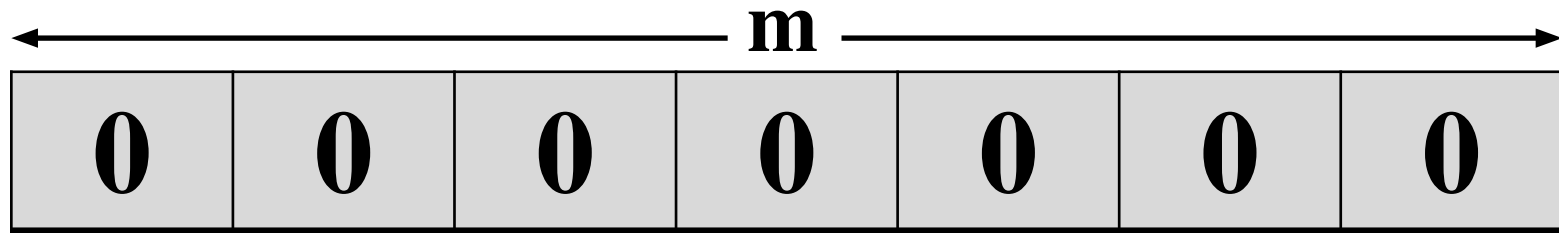
Dictionary

For most practical purposes:

$\epsilon = 2\%$, Bloom filter requires ≈ 8 bits/item

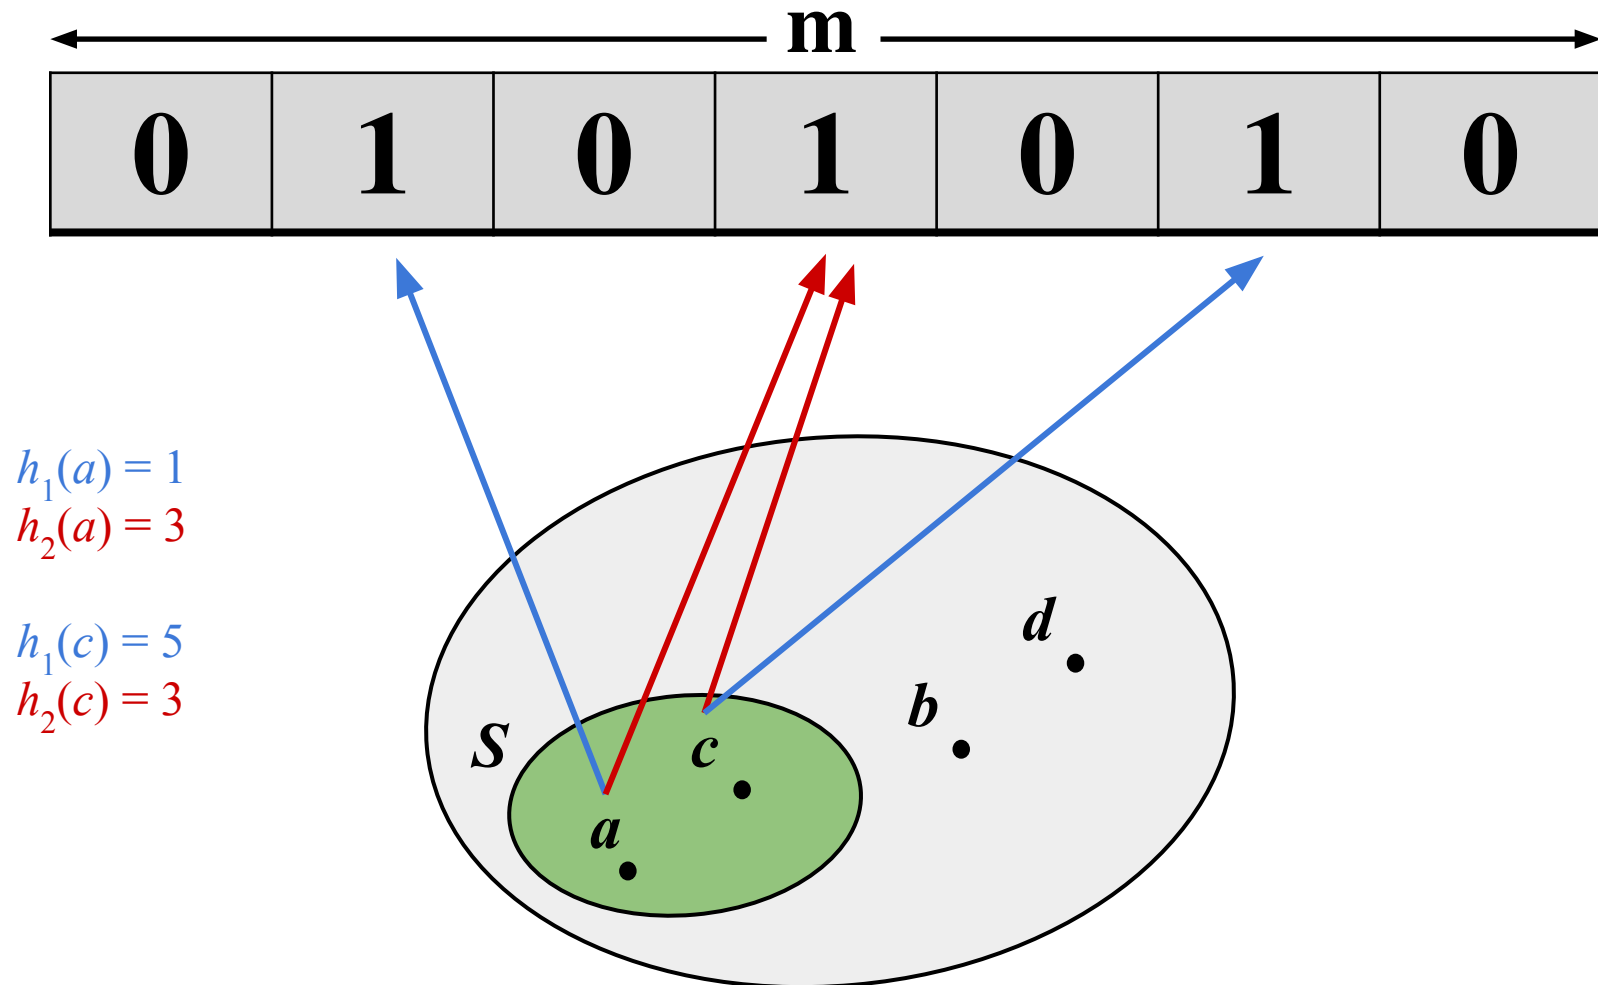
Classic filter: The Bloom filter [Bloom '70]

Bloom filter: a bit array + k hash functions




Classic filter: The Bloom filter [Bloom '70]

Bloom filter: a bit array + k hash functions (here $k = 2$)



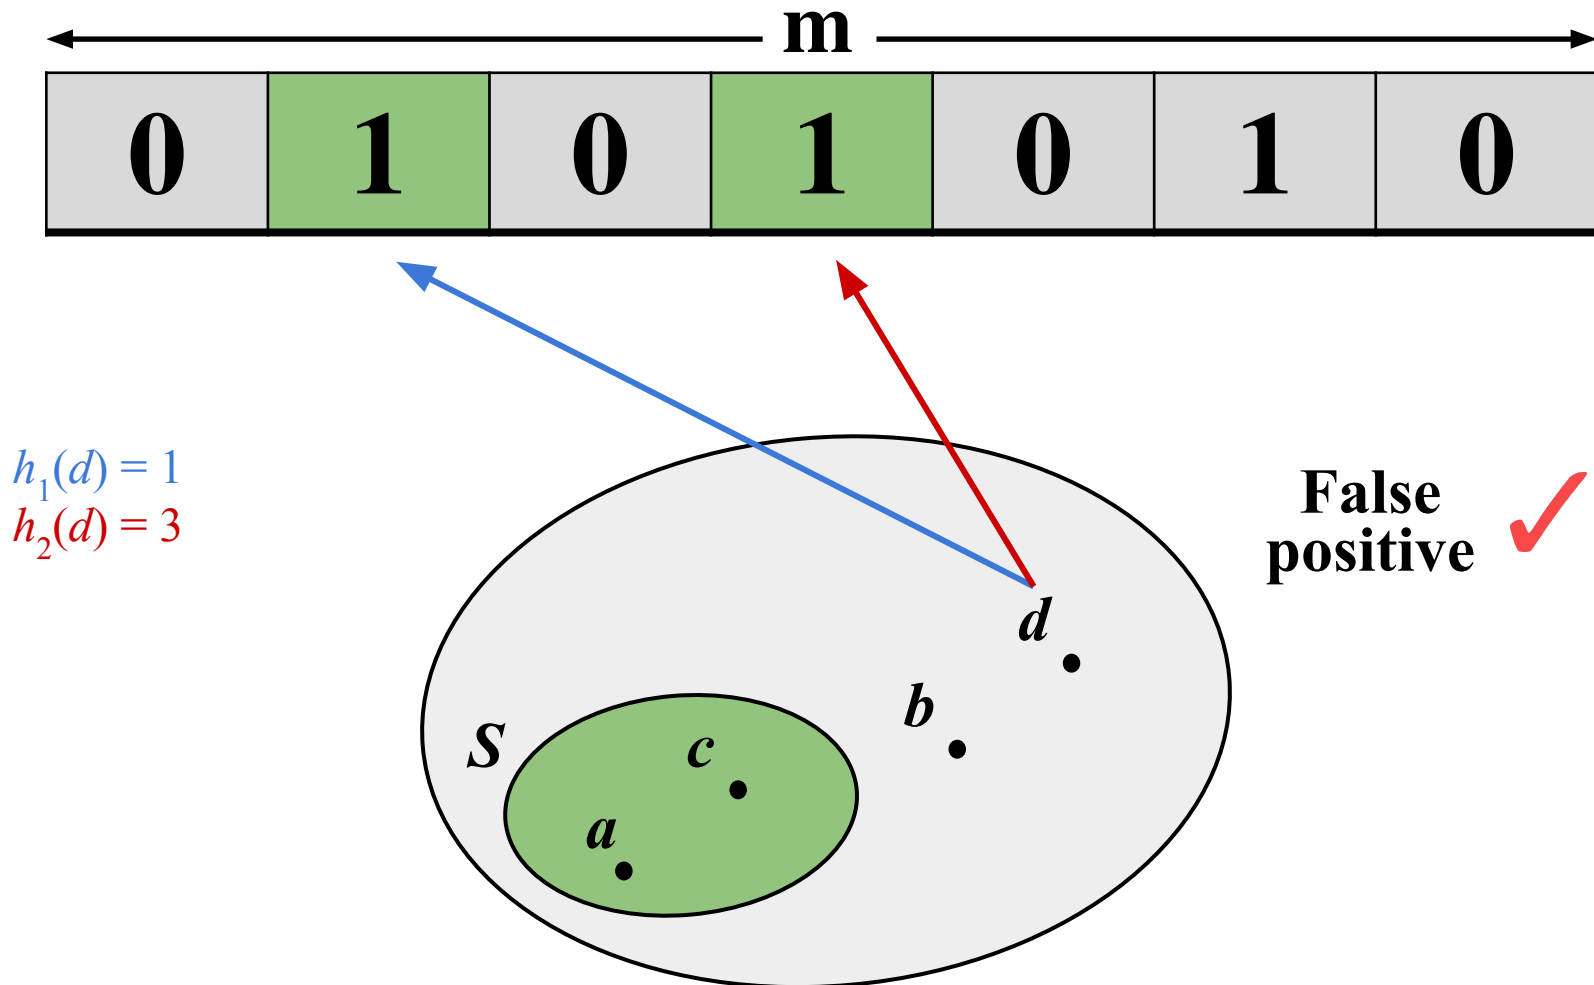
Classic filter: The Bloom filter [Bloom '70]

Bloom filter: a bit array + k hash functions (here $k=2$)



Classic filter: The Bloom filter [Bloom '70]

Bloom filter: a bit array + k hash functions (here $k=2$)



Bloom filter are ubiquitous (> 4300 citations)

Streaming applications



Networking



Databases



Computational biology



Storage systems



Bloom filter have suboptimal asymptotics

	Bloom filter	Optimal
Space	$\approx 1.44 n \log(1/\epsilon)$	$\approx n \log(1/\epsilon) + \Omega(n)$
CPU cost	$\Omega(1/\epsilon)$	$O(1)$
Data locality	$\Omega(1/\epsilon)$ probes	$O(1)$ probes

Application often work around Bloom filter limitations

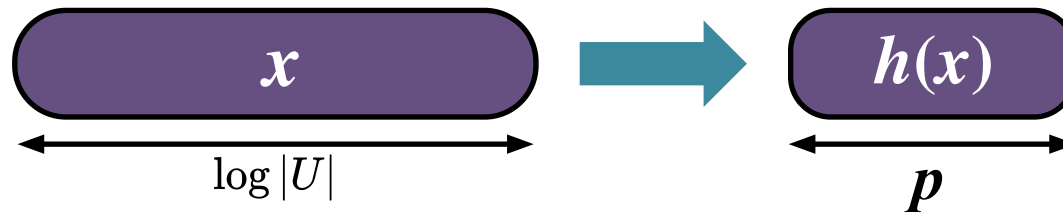
Limitations	Workarounds
No deletes	Rebuild
No resizes	Guess N , and rebuild if wrong
No filter merging or enumeration	???
No values associated with keys	Combine with another data structure

Bloom filter limitations increase system complexity, waste space, and slow down application performance

Quotienting is an alternative to Bloom filters

[Knuth. Searching and Sorting Vol. 3, '97]

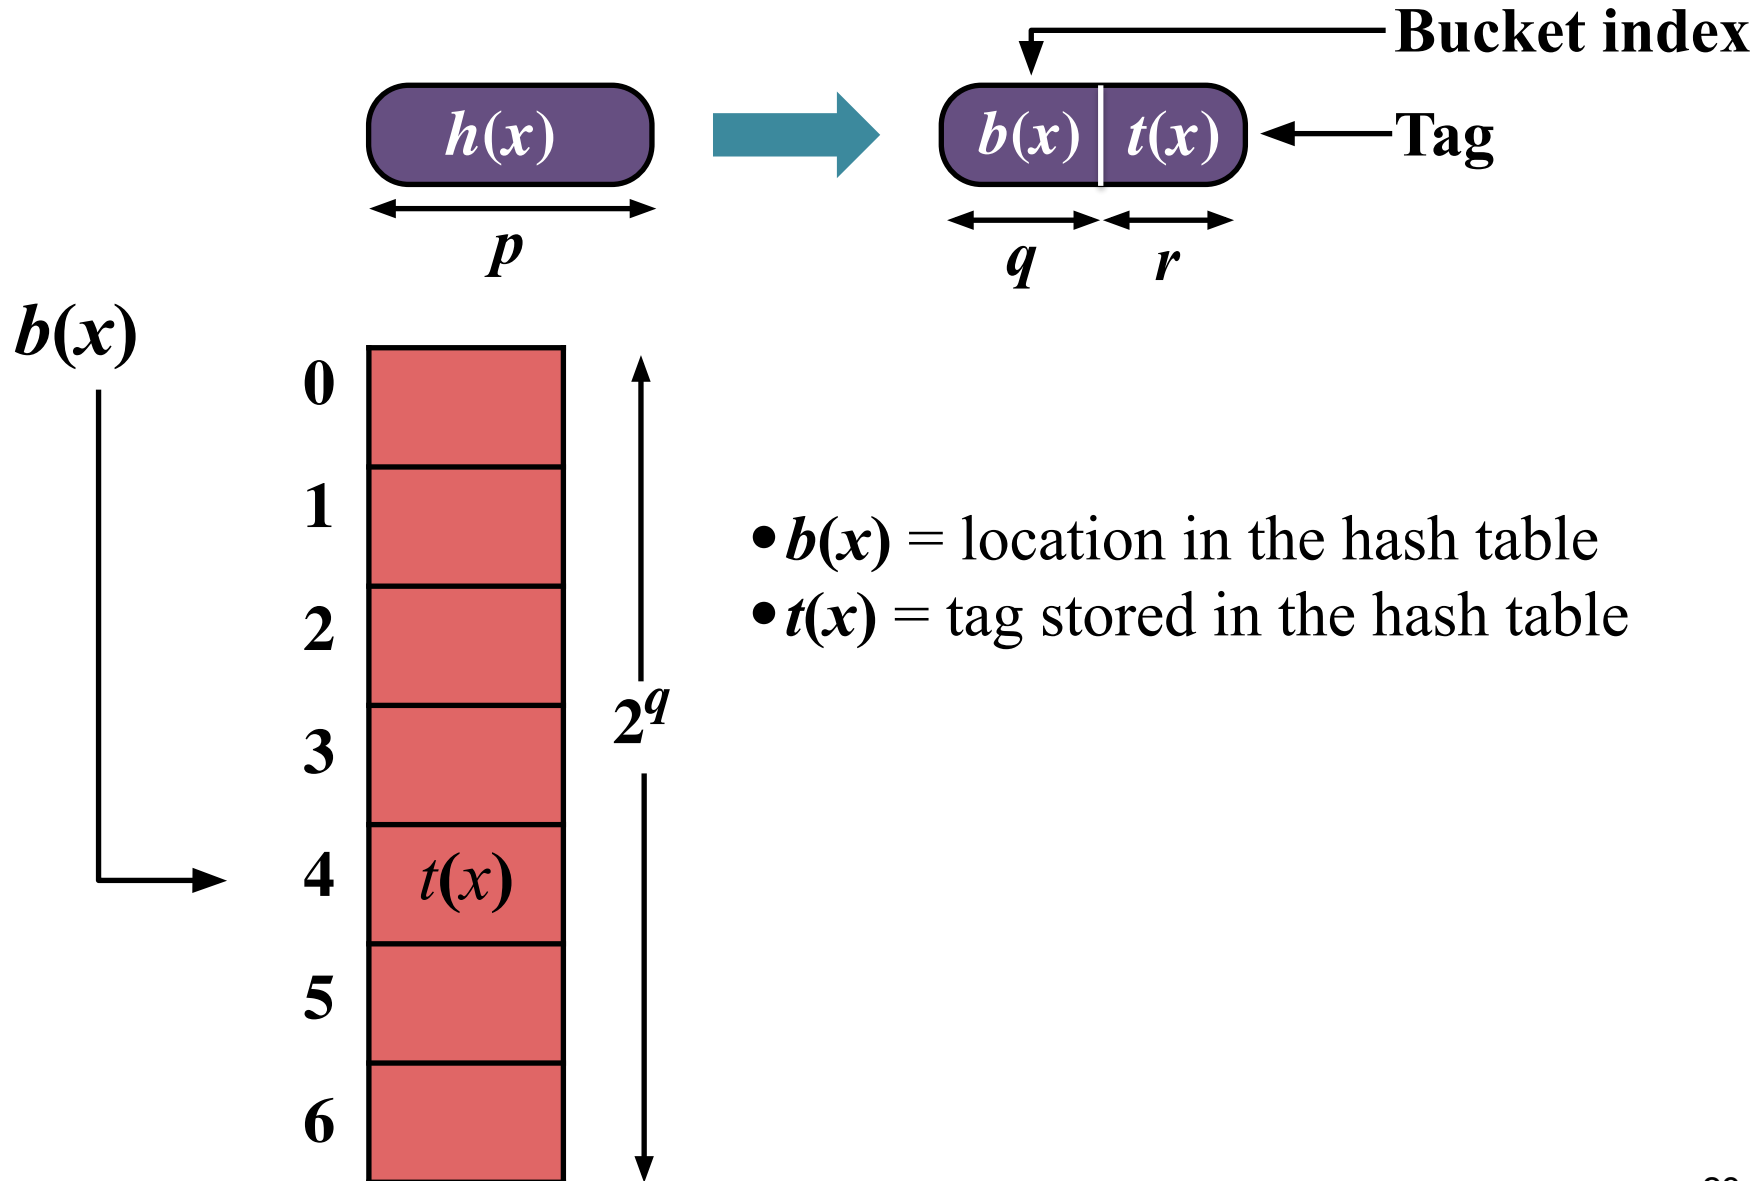
- **Store fingerprints compactly in a hash table.**
 - Take a fingerprint $h(x)$ for each element x .



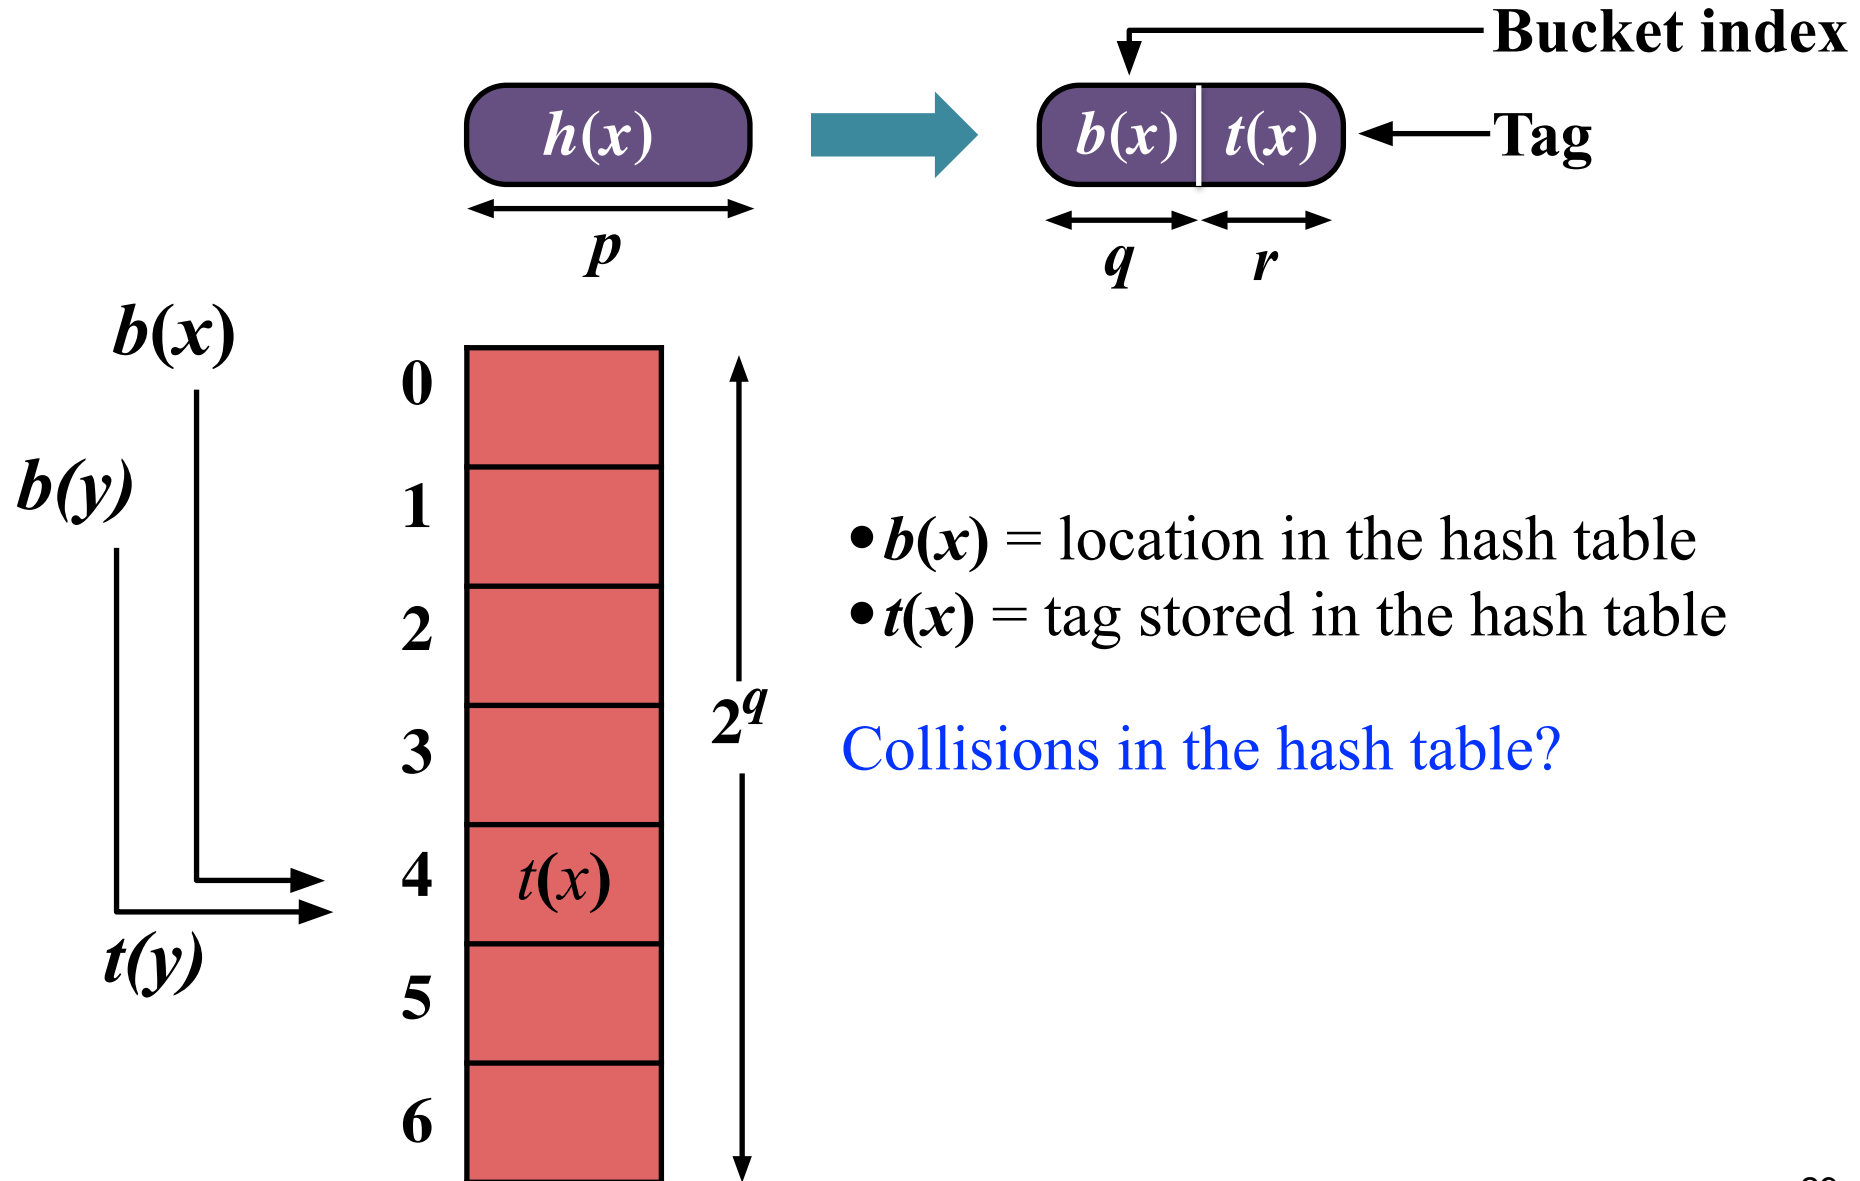
- **Only source of false positives:**
 - Two distinct elements x and y , where $h(x) = h(y)$
 - If x is stored and y isn't, $\text{query}(y)$ gives a false positives

$$\Pr[x \text{ and } y \text{ collide}] = \frac{1}{2^p}$$

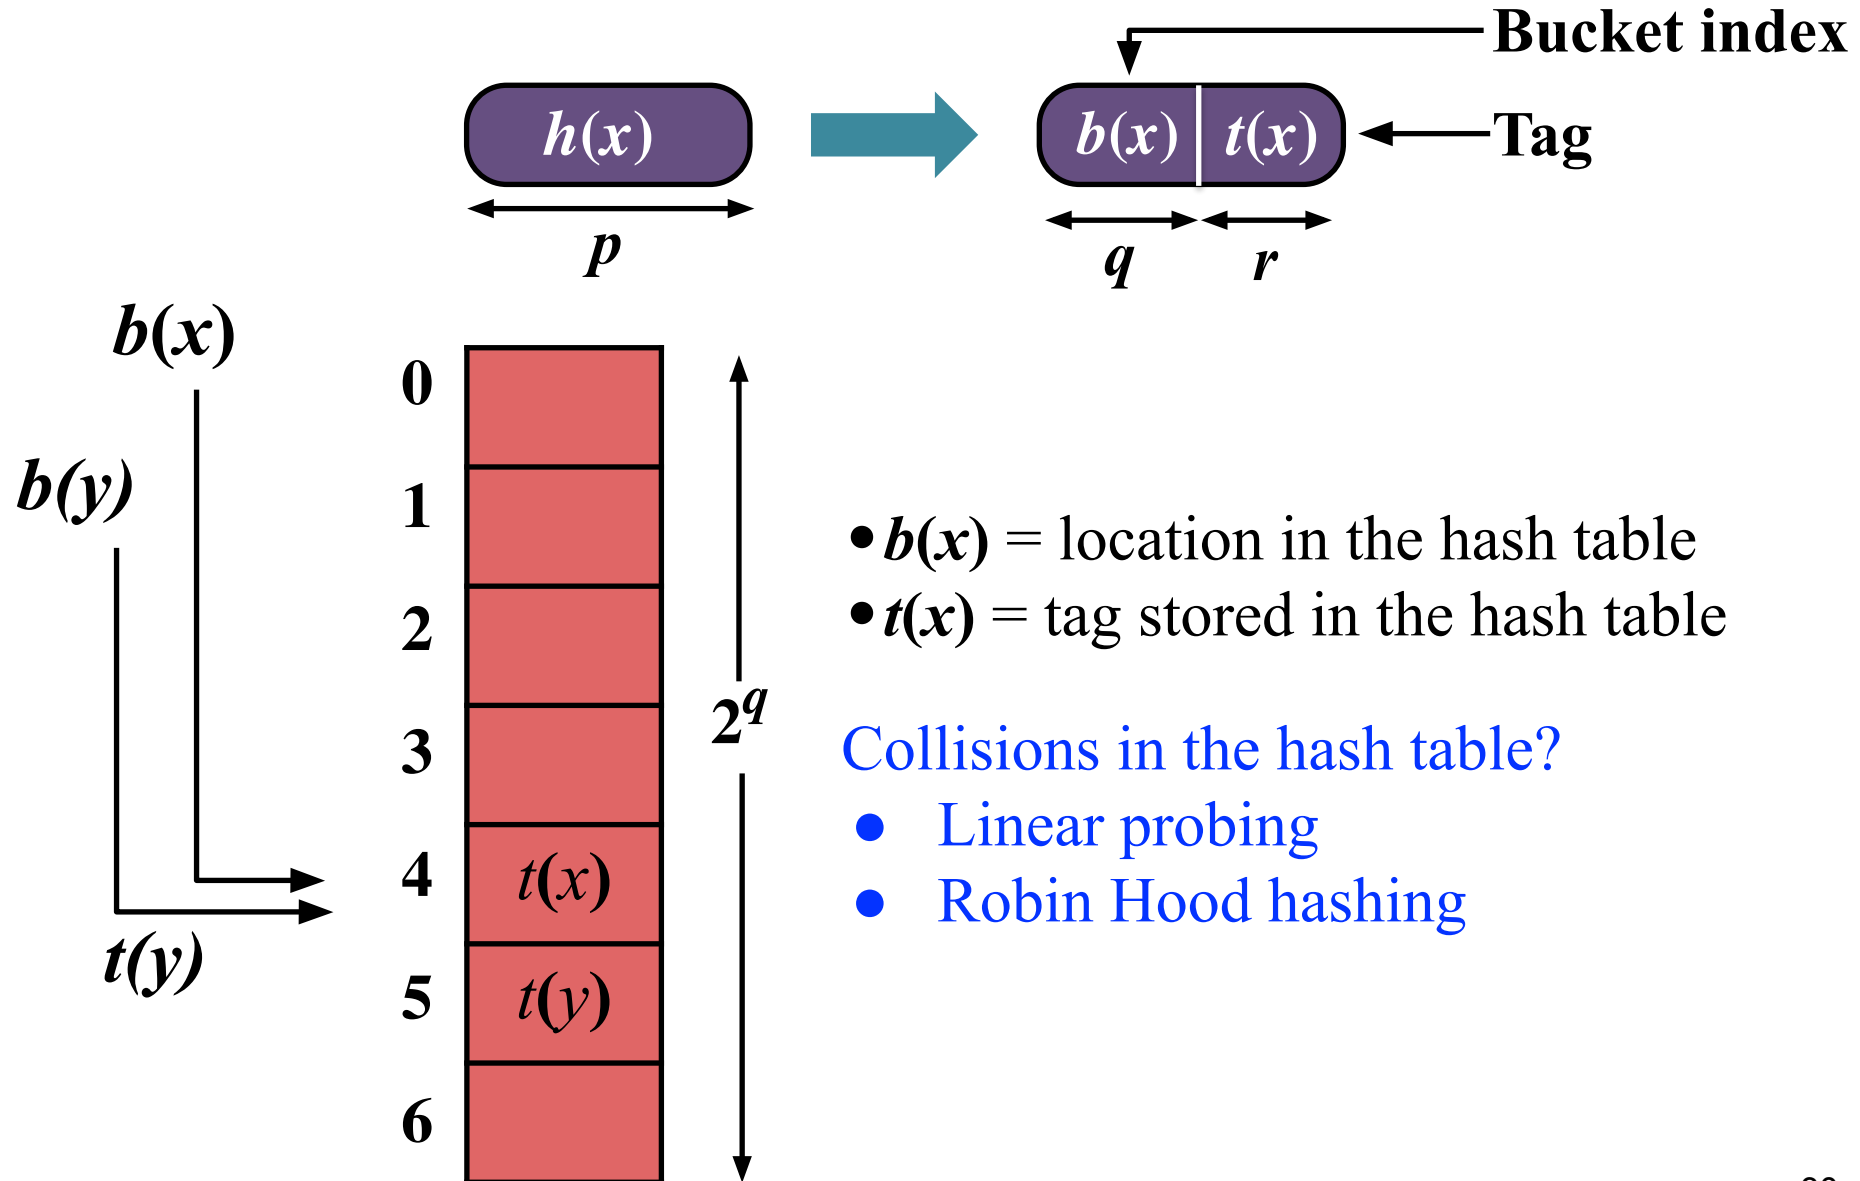
Storing fingerprints compactly



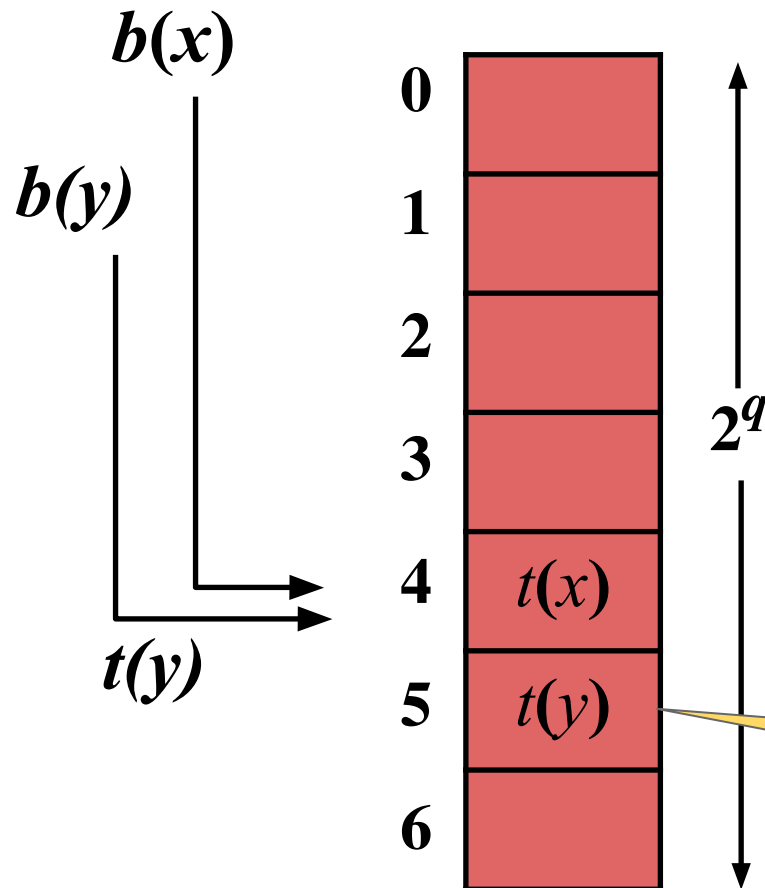
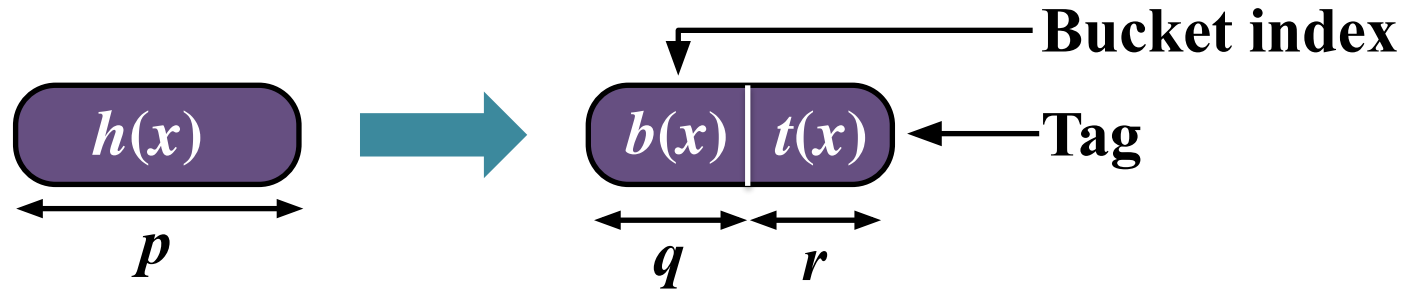
Storing fingerprints compactly



Storing fingerprints compactly



Storing fingerprints compactly



- $b(x)$ = location in the hash table
- $t(x)$ = tag stored in the hash table

Collisions in the hash table?

- Linear probing
- Robin Hood hashing

$t(y)$ belongs to slots 4 or 5?

Resolving collisions in the QF [Bender '12, Pandey '17]


- QF uses two metadata bits to resolve collisions and identify home bucket

	1		1				
	$t(u)$	$t(v)$	$t(w)$	$t(x)$	$t(y)$		

- The metadata bits group tags by their home bucket

Resolving collisions in the QF [Bender '12, Pandey '17]


- QF uses two metadata bits to resolve collisions and identify home bucket



- The metadata bits group tags by their home bucket

Resolving collisions in the QF [Bender '12, Pandey '17]

- QF uses two metadata bits to resolve collisions and identify home bucket



- The metadata bits group tags by their home bucket

The metadata bits enable us to identify the slots holding the contents of each bucket.

Quotienting enables many features in the QF

- Good cache locality
- Efficient scaling out-of-RAM
- Deletions
- Enumerability/Mergeability
- Resizing
- Maintains count estimates
- Uses variable-sized encoding for counts [Counting quotient filter]
 - **Asymptotically optimal space: $O(\sum |C(x)|)$**




Quotient filters use less space than Bloom filters for all practical configurations

	Quotient filter	Bloom filter	Optimal
Space	$\approx n \log(1/\epsilon) + 2.125n$	$\approx 1.44 n \log(1/\epsilon)$	$\approx n \log(1/\epsilon) + \Omega(n)$
CPU cost	$O(1)$ expected	$\Omega(1/\epsilon)$	$O(1)$
Data locality	1 probe + scan	$\Omega(1/\epsilon)$ probes	$O(1)$ probes

The quotient filter has theoretical advantages over the Bloom filter

Quotient filters use less space than Bloom filters for all practical configurations




Bloom filter: $\sim 1.44 \log(1/\epsilon)$ bits/element.

Quotient filter: $\sim 2.125 + \log(1/\epsilon)$ bits/element.


Quotient filters perform better (or similar) to other non-counting filters



Inserts



Lookups



- Insert performance is similar to the state-of-the-art non-counting filters
- Query performance is significantly fast at low load-factors and slightly slower at higher load-factors

Quotient filter's impact in computer science

Computational biology

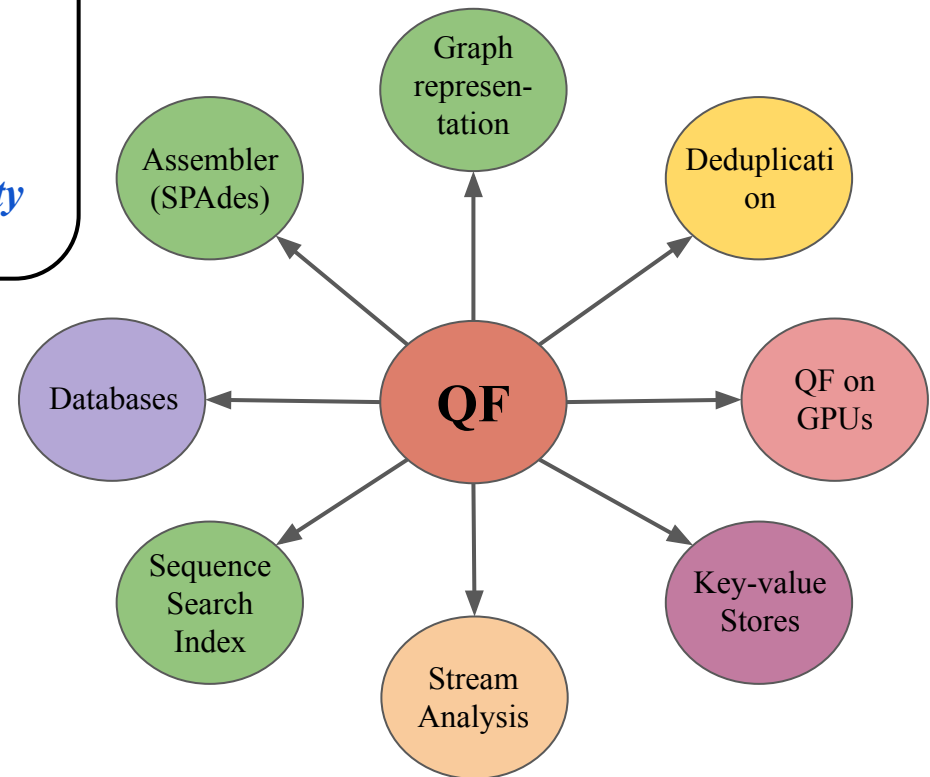
1. Squeakr
2. deBGR
3. Mantis
4. SPAdes assembler
5. Khmer software
6. MQF
7. VariantStore

Industry

1. VMware
2. Nutanix
3. Apocrypha
4. Hyrise
5. *A data security startup*

Databases/Systems

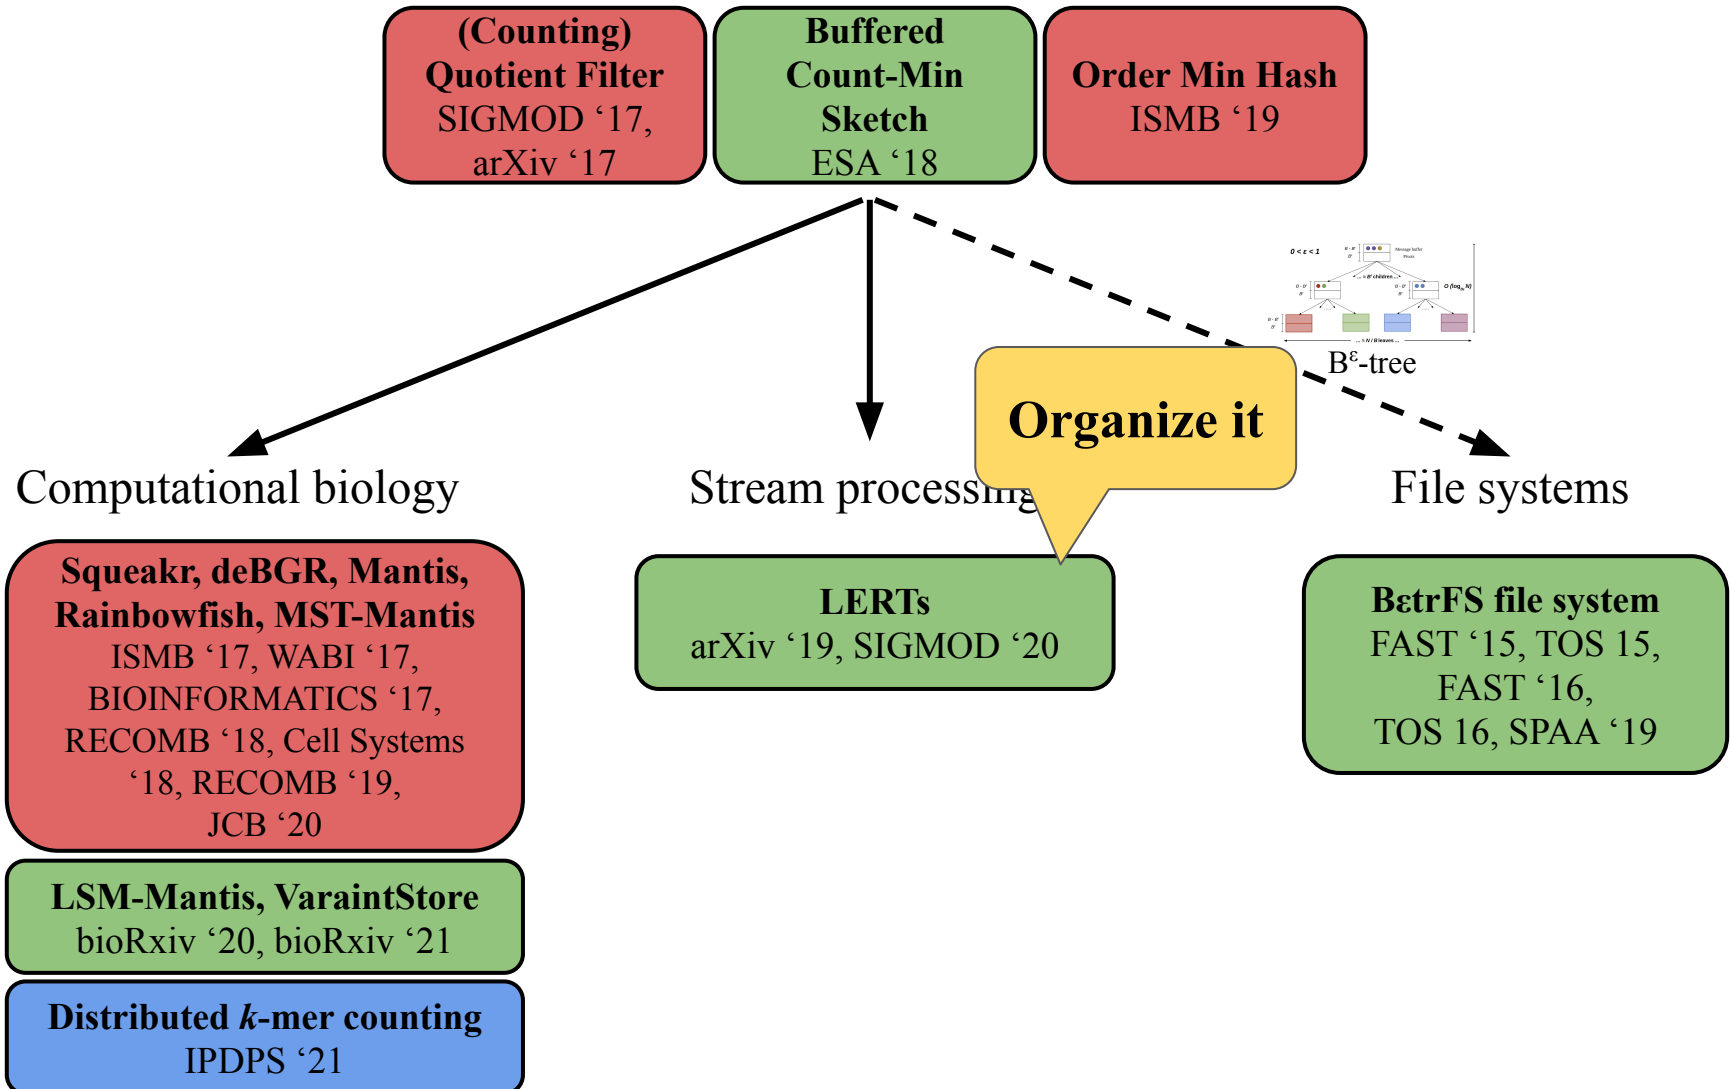
1. Counting on GPUs
2. Concurrent filters
3. Anomaly detection
4. BetrFS file system



Theoretically well-founded data structures can have a **big impact** on multiple subfields across *academia and industry*

Learned “Shrink it”. Now “Organize it”

Data structures & Algorithms



Open problem in stream processing

- A **high-speed stream** of key-value pairs arriving over time
- **Goal:** report every key **as soon as** it appears **T times** without missing any
- Firehose benchmark (Sandia National Lab) simulates the stream
<https://firehose.sandia.gov/>



Why should we care about this problem

Defense systems for cyber security

monitor high-speed streams for
malicious traffic

➡ **Catch all malicious events**

Malicious traffic forms a small
portion of the stream


➡ **Small reporting threshold**

Automated systems take defensive
actions for every reported event

➡ **Minimize false positives**


Timely event detection problem

- Stream of elements arrive over time




Timely event detection problem

- Stream of elements arrive over time
- An **event** occurs at time t if S_t occurs exactly T times in (s_1, s_2, \dots, s_t)




Timely event detection problem

- Stream of elements arrive over time
- An **event** occurs at time t if S_t occurs exactly T times in (s_1, s_2, \dots, s_t)



Timely event detection problem

- Stream of elements arrive over time
- An **event** occurs at time t if S_t occurs exactly T times in (s_1, s_2, \dots, s_t)
- In **timely event-detection problem (TED)**, we want to report all events shortly after they occur.



Features we need in the solution

- Stream is large (e.g., terabytes) and high-speed (millions/sec)

High throughput ingestion



Features we need in the solution

- Stream is large (e.g., terabytes) and high-speed (millions/sec)

High throughput ingestion

- Events are high-consequence real-life events

No false-negatives; few false-positives

Timely reporting (real-time)



Features we need in the solution

- Stream is large (e.g., terabytes) and high-speed (millions/sec)

High throughput ingestion

- Events are high-consequence real-life events

No false-negatives; few false-positives

Timely reporting (real-time)


- Malicious traffic forms a small portion of the stream

Very small reporting thresholds




One-pass streaming has errors

- **Heavy hitter problem:** report items whose frequency $\geq \phi N$
- Exact one-pass solution requires $\Omega(N)$ space



One-pass streaming has errors

- **Approximate solution:** report all items with count $\geq \phi N$, none with $< (\phi - \epsilon)N$ [Alon et al. 96, Berinde et al. 10, Bhattacharyya et al. 16, Bose et al. 03, Braverman et al. 16, Charikar et al. 02, Cormode et al. 05, Demaine et al. 02, Dimitropoulos et al. 08, Larsen et al. 16, Manku et al. 02.]
- Approximate solutions requires: $\Omega(1/\epsilon)$



Real time with false-positives!


One-pass streaming has errors

- **Approximate solution:** report all items with count $\geq \phi N$, none with $< (\phi - \epsilon)N$ [Alon et al. 96, Berinde et al. 10, Bhattacharyya et al. 16, Bose et al. 03, Braverman et al. 16, Charikar et al. 02, Cormode et al. 05, Demaine et al. 02, Dimitropoulos et al. 08, Larsen et al. 16, Manku et al. 02.]
- Approximate solutions requires: $\Omega(1/\epsilon)$

For Sandia, ϕN is a small constant (e.g., 24),
So $\Omega(1/\epsilon)$ is very very large!!
Can't solve in RAM for very small ϕ

RAM
[Misra & Gries '82]

RAM



Real time with false-positives!

One-pass solution has:

- Stream is large (e.g., terabytes) and high-speed (millions/sec)

High throughput ingestion



- Events are high-consequence real-life events

No false-negatives; few false-positives



Timely reporting (real-time)



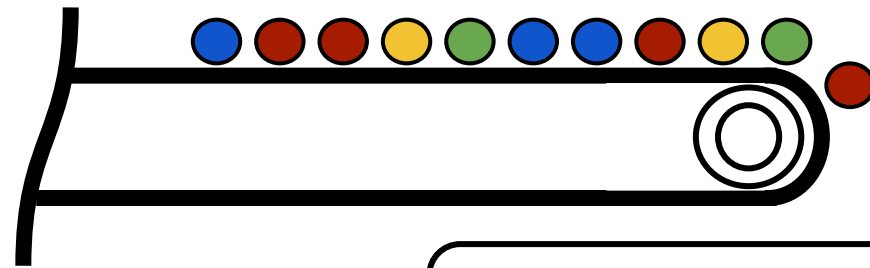
- Malicious traffic forms a small portion of the stream

Very small reporting thresholds



Two-pass streaming isn't real-time

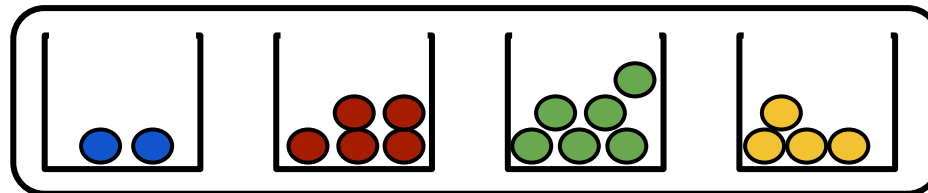
- A second pass over the stream can get rid of errors
- Store the stream on SSD and access it later



Scales to very small ϕ
but offline!

SSD

RAM



Second pass



Two-pass solution has:

- Stream is large (e.g., terabytes) and high-speed (millions/sec)

High throughput ingestion



- Events are high-consequence real-life events

No false-negatives; few false-positives



Timely reporting (real-time)



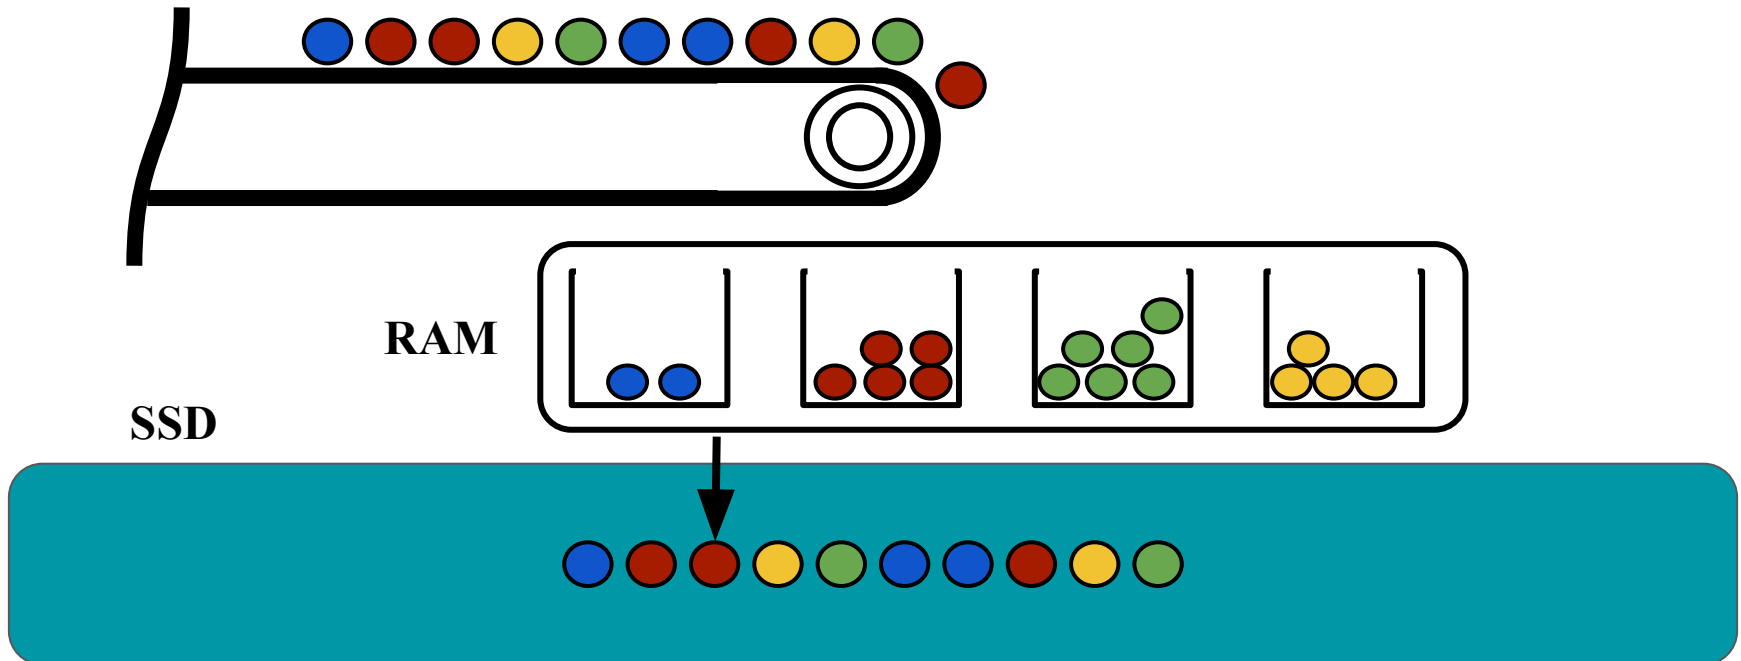
- Malicious traffic forms a small portion of the stream

Very small reporting thresholds



If data is stored: why not access it?

Why wait for second pass?





Idea: combine Streaming and EM

Use an efficient external-memory counting data structure to scale Misra-Gries algorithm to SSDs

**Streaming
model**



**External memory
algorithms**


External memory model [Aggarwal+Vitter '08]

- **How computations work:**

- Data is transferred in blocks between RAM and disk.
- The number of block transfers dominate the running time.


- **Goal: Minimize number of block transfers**

- Performance bounds are parameterized by block size B , memory size M , data size N .



Cascade filter: write-optimized quotient filter

[Bender et al. '12, Pandey et al. '17]



- The Cascade filter efficiently scales out-of-RAM
- It accelerates insertions at some cost to queries

Cascade filter operations

Insert	Query
$O\left(\frac{1}{B} \log \frac{N}{M}\right)$	$O\left(\log \frac{N}{M}\right)$

Cascade filter operations

Insert	Query
$O\left(\frac{1}{B} \log \frac{N}{M}\right)$	$O\left(\log \frac{N}{M}\right)$

< 1 I/O per
observation



Cascade filter operations

Insert	Query
$O\left(\frac{1}{B} \log \frac{N}{M}\right)$	$O\left(\log \frac{N}{M}\right)$

< 1 I/O per
observation



> 1 I/O per
observation



Cascade filter doesn't have real-time reporting

But every insert is also a query in real-time reporting!

Insert	Query
$O\left(\frac{1}{B} \log \frac{N}{M}\right)$	$O\left(\log \frac{N}{M}\right)$

< 1 I/O per observation



> 1 I/O per observation



Cascade filter doesn't have real-time reporting

But every insert is also a query in real-time reporting!

Insert	Query
$O\left(\frac{1}{B} \log \frac{N}{M}\right)$	$O\left(\log \frac{N}{M}\right)$

Traditional cascade filter doesn't solve the problem!

observation 


observation 



Idea: reporting with bounded delay

We define the time stretch of a report to be

$$\text{Time stretch} = 1 + \alpha = 1 + \frac{\text{Delay}}{\text{Lifetime}}$$






Idea: reporting with bounded delay

We define the time stretch of a report to be

$$\text{Time stretch} = 1 + \alpha = 1 + \frac{\text{Delay}}{\text{Lifetime}}$$

Main idea: the longer the lifetime of an item, the more leeway we have in reporting it



Leveled External-Memory Reporting Table (LERT) [Pandey '20]

- Given a stream of size N and $\phi N > \Omega(N/M)$ the amortized cost of solving real-time event detection is

$$O \left(\left(\frac{1}{B} + \frac{1}{(\phi - 1/M)N} \right) \log \frac{N}{M} \right)$$

- For a **constant** α , can support arbitrarily small thresholds ϕ with amortized cost

$$O \left(\frac{1}{B} \log \frac{N}{M} \right)$$

Takeaway: Online reporting comes at the cost of throughput but almost online reporting is essentially free!

Leveled External-Memory Reporting Table (LERT) [Pandey '20]

- Given a stream of size N and $\phi N > \Omega(N/M)$ the amortized cost of solving real-time event detection is

Can achieve timely reporting at effectively the optimal insert cost; no query cost

with amortized cost


$$O\left(\frac{1}{B} \log \frac{N}{M}\right)$$

Takeaway: Online reporting comes at the cost of throughput but almost online reporting is essentially free!

Evaluation


- Empirical timeliness
- High-throughput ingestion

Evaluation: empirical time stretch



Average time stretch is 43% smaller than theoretical upper bound.

Evaluation: scalability



The insertion throughput increases as we add more threads.
We can achieve $> 13\text{M}$ insertions/sec.

LERT: supports scalable and real-time reporting

- Stream is large (e.g., terabytes) and high-speed (millions/sec)

High throughput ingestion

- Events are high-consequence real-life events

No false-negatives; few false-positives

Timely reporting (real-time)

- Malicious traffic forms a small portion of the stream

Very small reporting thresholds



Future work overview



Data Science

Scalable Data Systems


Data structures &
Algorithms

Future work: Data Structures & Algorithms

Goal: Overcome *decades-old* data structure *trade-offs* using modern hardware and new algorithmic paradigms


Trade-off 1: Insertion throughput degrades with load factor

Insertion throughput vs load factor of state-of-the-art filters




Many update-intensive applications (e.g., network caches, data analytics, etc.) maintain filters at high load factors

Trade-off 1: Insertion throughput degrades with load factor




Performance suffers due to high-overhead of *collision resolution*

Combining techniques + new hardware



Combining hashing techniques (**Robin Hood + 2-choice hashing**)
Using ultra-wide vector operations (**AVX512-BW**)

Combining techniques + new hardware

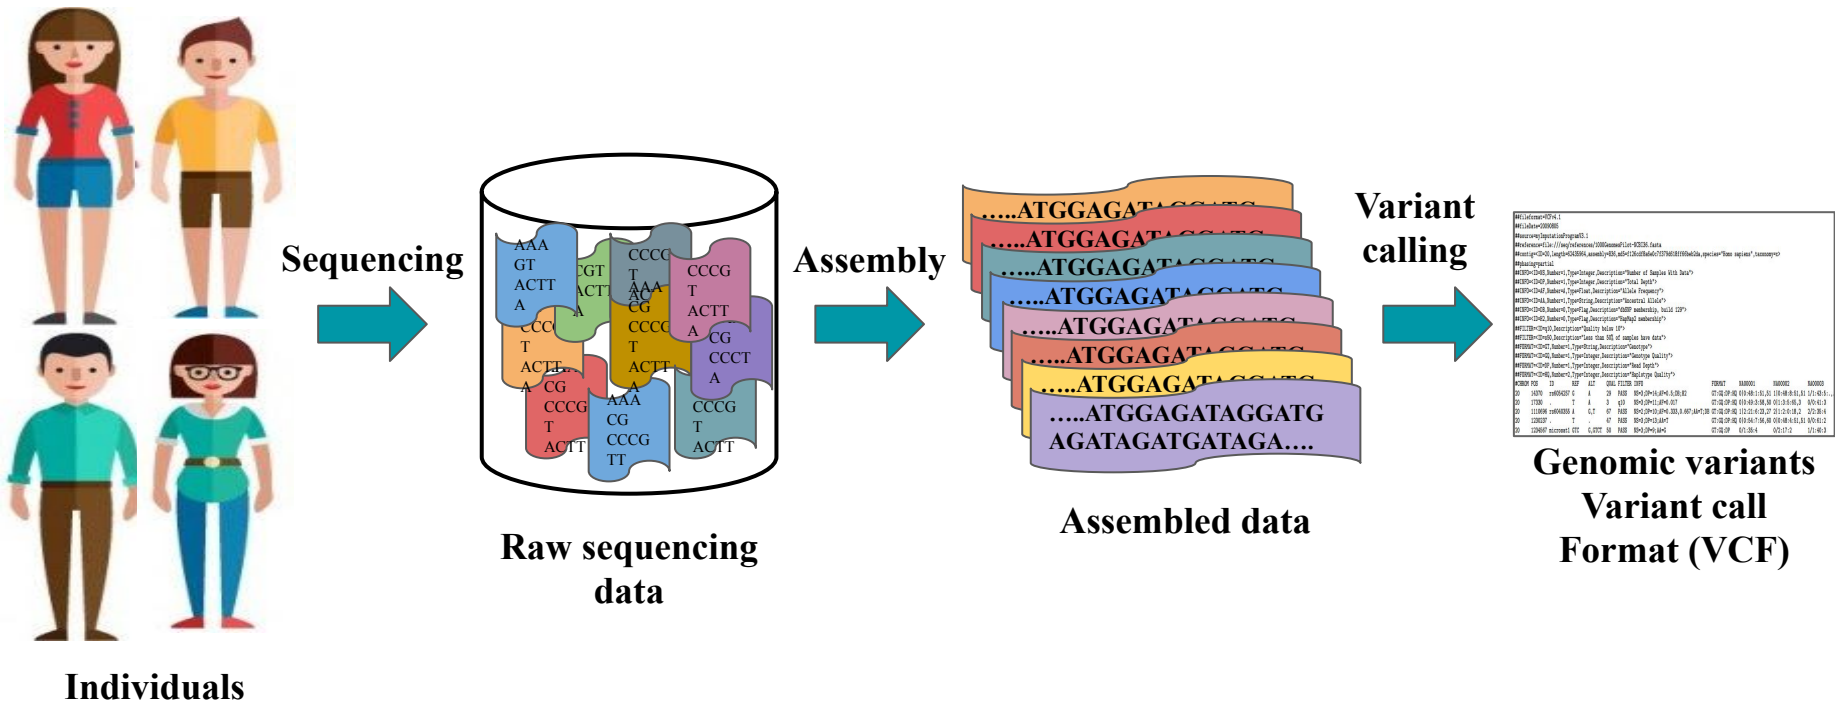


Combining hashing techniques (**Robin Hood + 2-choice hashing**)
Using ultra-wide vector operations (**AVX512-BW**)

Future work: Data Systems

Goal: Build a *population-scale* index on variation data to enable downstream apps gain *quick insights into variants*

Country-scale sequencing efforts produce huge amounts of sequencing data



- 1000 Genomes project [<https://www.internationalgenome.org/>]
- The Cancer Genome Atlas (TCGA) [<https://portal.gdc.cancer.gov/>]
- Genotype-Tissue Expression (GTEx) [<https://gtexportal.org/home/>]


Variation data analysis can improve downstream applications

- Population-level disease analysis
- Genome-wide association studies
- Personalized medicine
- Cancer remission-rate prediction
- Colocalization analysis
- PCR primer design
- Genome assembly



Individuals

Sequencing & assembly



Population Genomes

For person P , return the closest variant from position X

Count the number of variants in a gene

List all people, with $> N$ variants in a gene

Return all positions with variants in a gene

List all people, with sequence S in a gene

Indexing in multiple coordinates is challenging

Reference-only indexes map positions only in the reference coordinate system

$$f(p_i, p_j) \rightarrow (v_i \dots v_n), \text{ where } p_i \leq p_j$$

Pan-genome analysis involves queries based on sample coordinate systems

$$\begin{array}{l} \text{Num} \\ \text{Samples} \end{array} \left\{ \begin{array}{l} f_1(p_i, p_j) \rightarrow (v_i \dots v_n), \text{ where } p_i \leq p_j \\ \vdots \\ f_s(p_i, p_j) \rightarrow (v_i \dots v_n), \text{ where } p_i \leq p_j \end{array} \right.$$

Maintaining thousands of mappings ***increases*** computational ***complexity*** and ***memory footprint***

Limits scalability to population-scale data

Indexing in multiple coordinates is challenging

Reference-only indexes map positions only in the reference coordinate system

$$f(p_i, p_j) \rightarrow (v_i \dots v_n), \text{ where } p_i \leq p_j$$

Pan-genome analysis involves queries based on sample coordinate systems

Existing systems don't support multiple coordinate systems. The ones that do, don't *scale* beyond a few thousand samples.

$$f_s(p_i, p_j) \rightarrow (v_i \dots v_n), \text{ where } p_i \leq p_j$$

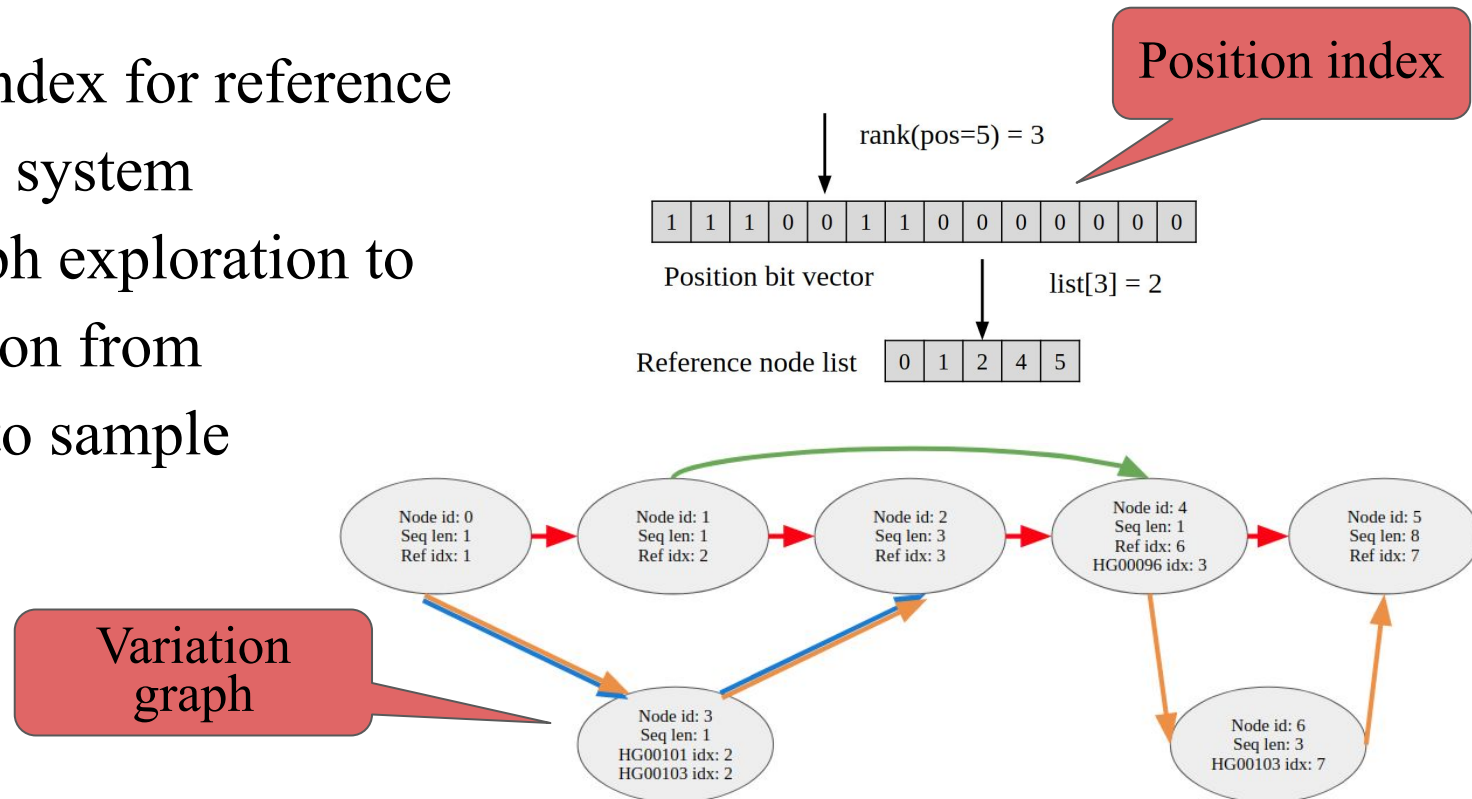
Maintaining thousands of mappings *increases* computational *complexity* and *memory footprint*

Limits scalability to population-scale data

An inverted index on the pan-genome graph

- Partition the variation graph based on coordinate ranges
 - Store partitions on disk
- } Queries often require loading 1-2 partitions

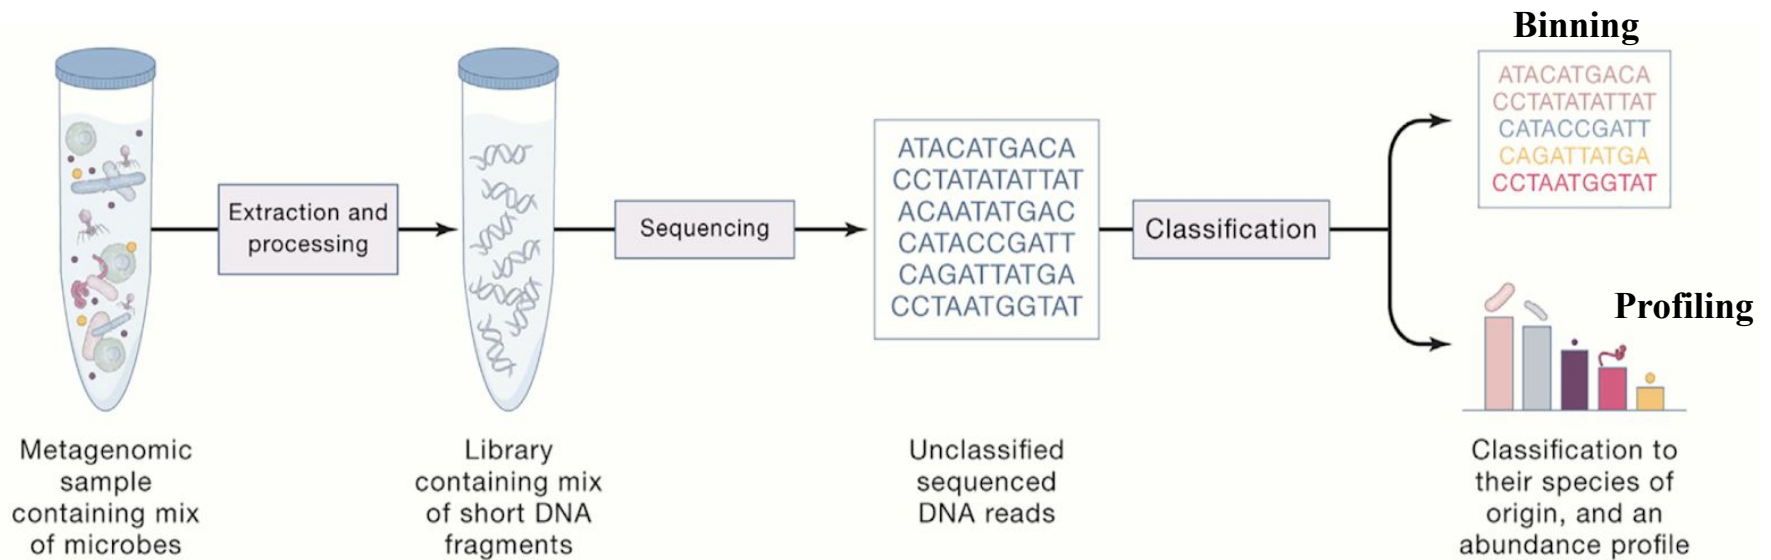
- Succinct index for reference coordinate system
- Local-graph exploration to map position from reference to sample coordinate



Future work: Data Science for genomics


Goal: Classification of metagenomic reads and *identification of novel species* using *graph neural networks* (GNN)

Metagenomic classification pipeline



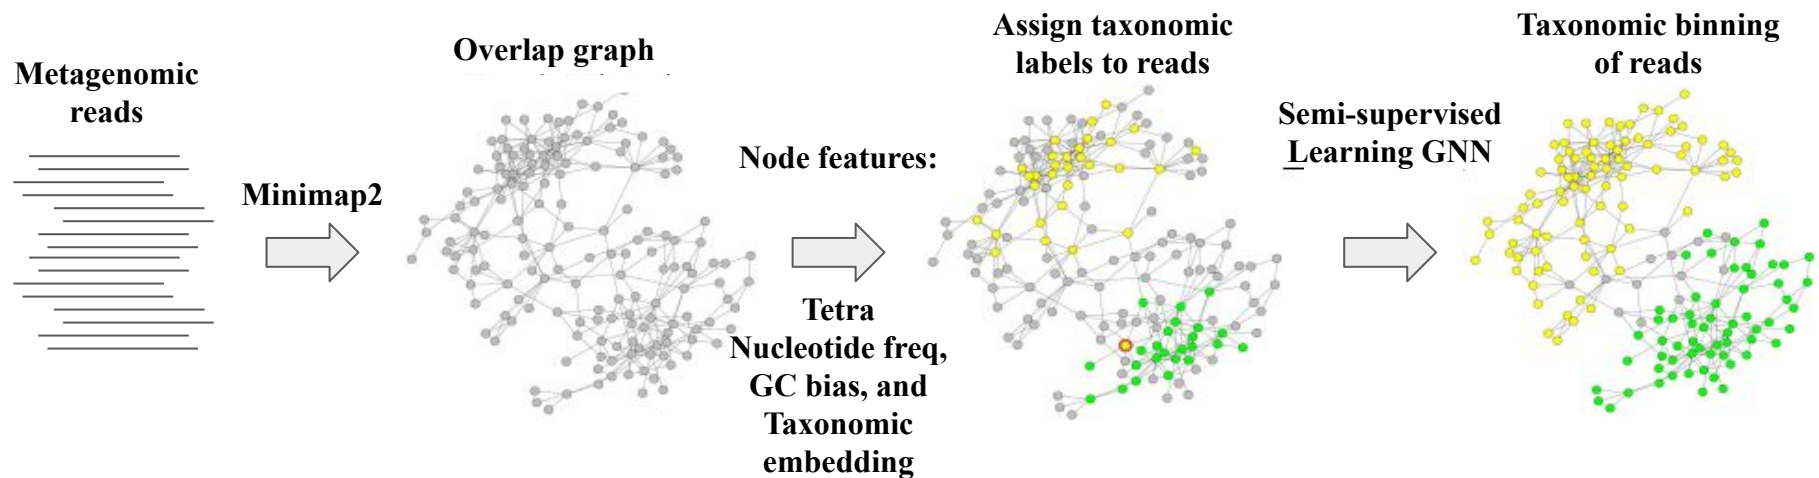
[Ye et al. 2019]

Existing techniques offer low recall



Classification is done based *only on the read contents*


Use overlap relationship between reads



- Generate overlap graph: reads → nodes & overlap → edges
- Node features → Tetra nucleotide freq of reads
- Reference-based mapping as ground truth labels

Overlap graph + graph neural network (GNN)

Classification accuracy



Can achieve high recall using graph learning

Conclusion

- Scalability of data management systems will be the biggest challenge in future
- Changing hardware give rise to new algorithmic paradigms

Data Science at Scale

ML

Genomics

Cyber Sec.

NLP

Data Systems

Data structures & Algorithms

Scale down

Scale to disk

Scale out

Modern hardware

Vector inst.

GPU

NVM

SSD

We need to *redesign* existing data structures to take full advantage of modern hardware and *rebuild* data systems to efficiently support *future* data science.

<https://prashantpandey.github.io>



Backup slides

Quotient filter design

Implementation:

2 Meta-bits per slot.

$h(x) \rightarrow h_0(x) \parallel h_1(x)$





Quotient filter design

Implementation:

2 Meta-bits per slot.

$h(x) \rightarrow h_0(x) \parallel h_1(x)$





Quotient filter design

Implementation:

2 Meta-bits per slot.

$h(x) \rightarrow h_0(x) \parallel h_1(x)$





Quotient filter design

Implementation:

2 Meta-bits per slot.

$h(x) \rightarrow h_0(x) \parallel h_1(x)$

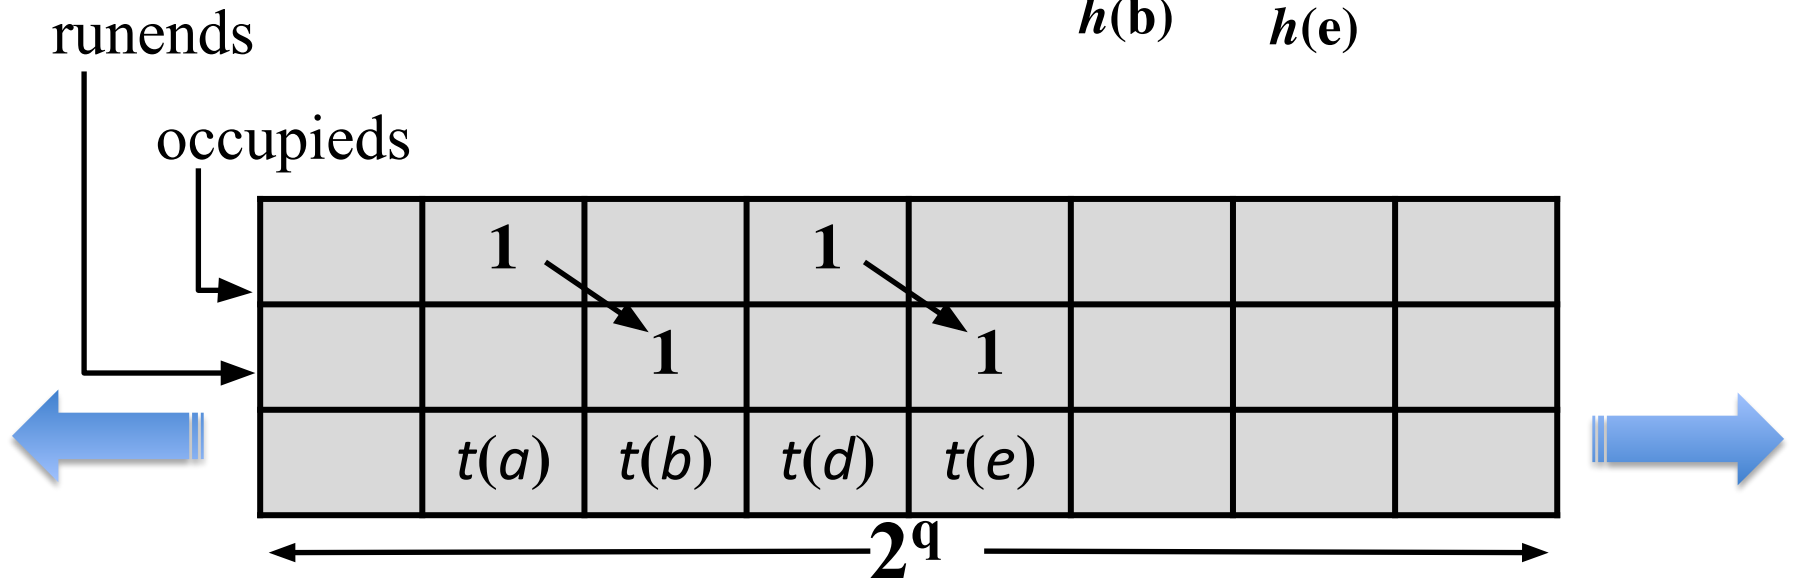



Quotient filter design

Implementation:

2 Meta-bits per slot.

$h(x) \rightarrow h_0(x) \parallel h_1(x)$





Quotient filter design

Implementation:

2 Meta-bits per slot.


$h(x) \rightarrow h_0(x) \parallel h_1(x)$



Back

Quotient filters can also be exact

- Quotient filters store $h(x)$ exactly
- To store x exactly, use an invertible hash function

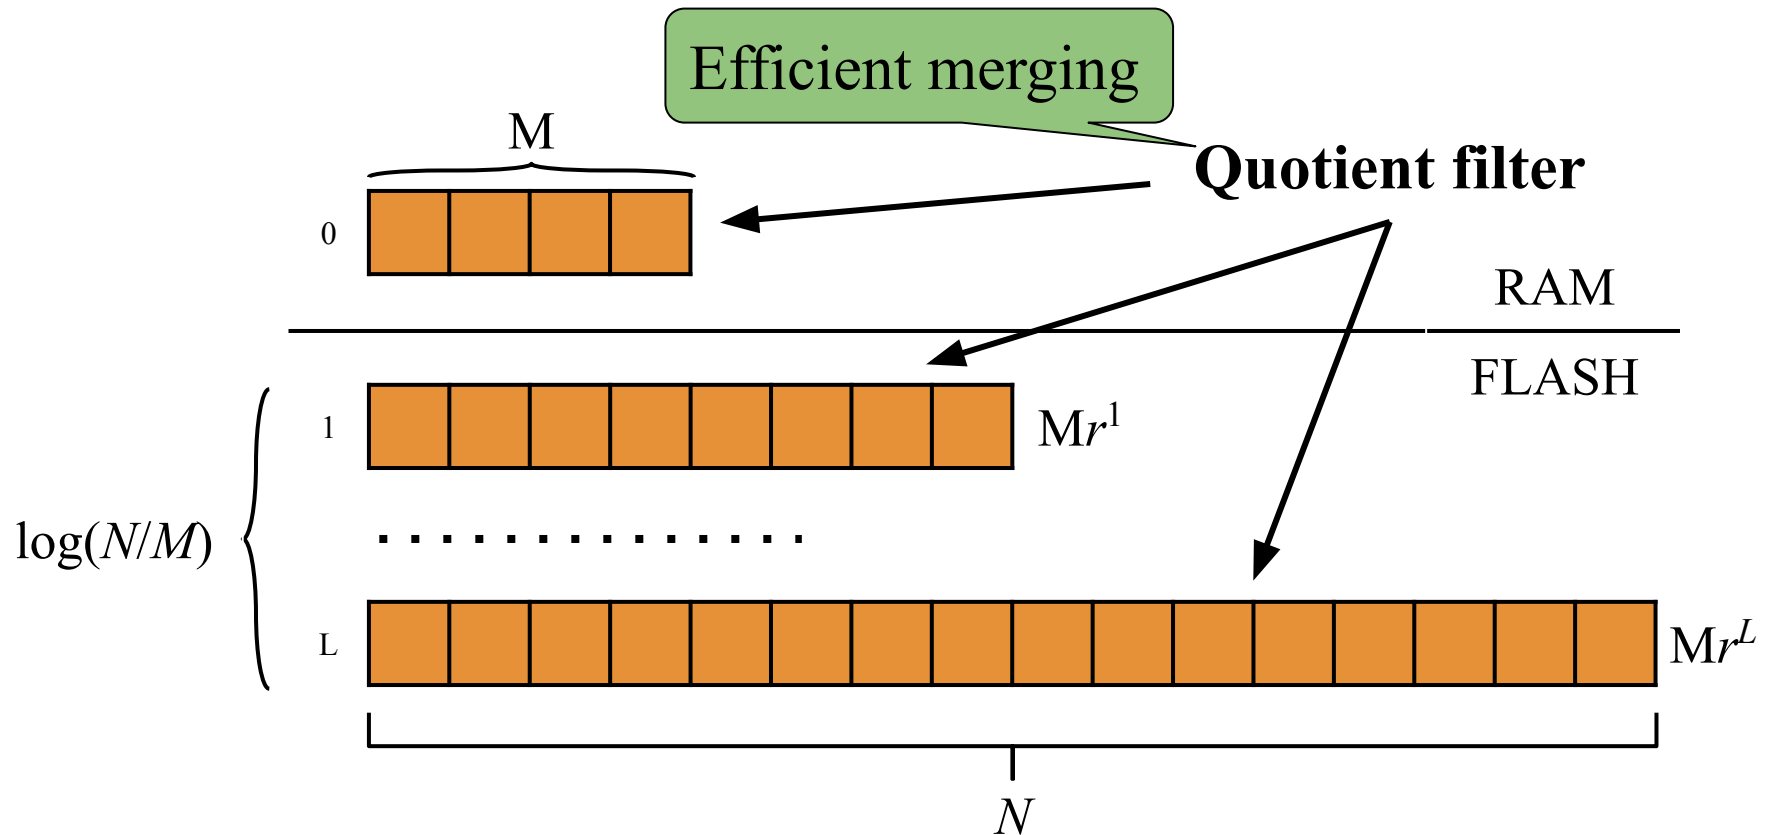


- For n elements and p -bit hash function:

Space usage: $\sim p \log_2 n$ bits/element

Cascade filter: write-optimized quotient filter

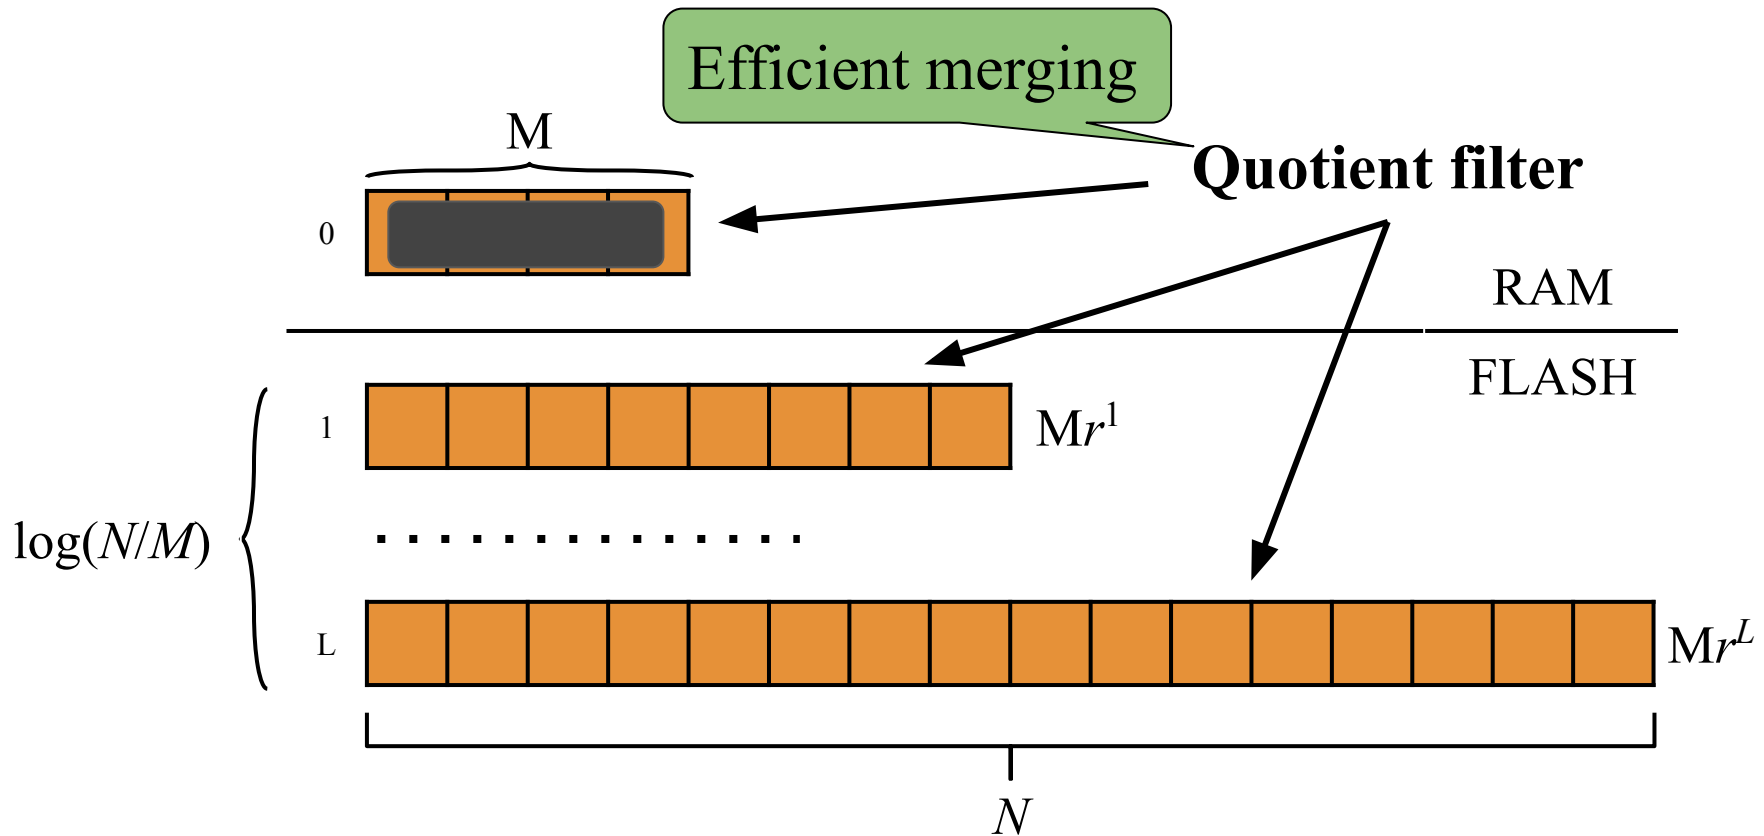
[Bender et al. '12, Pandey et al. '17]



- The Cascade filter efficiently scales out-of-RAM
- It accelerates insertions at some cost to queries

Cascade filter: flushing

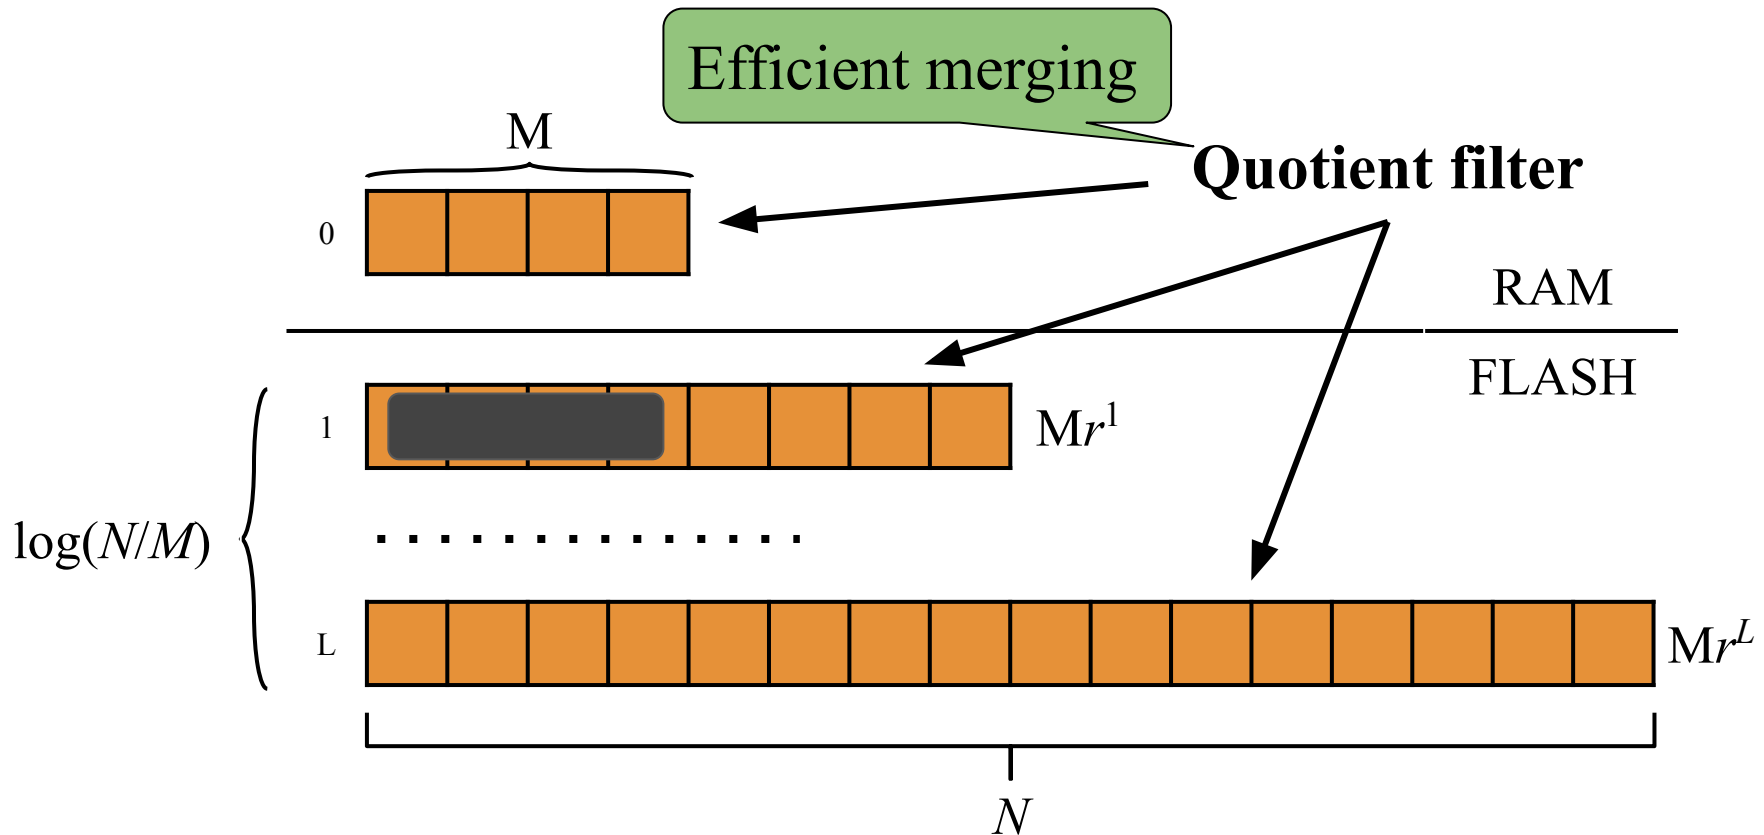
[Bender et al. '12, Pandey et al. '17]



Items are initially inserted in the RAM level

Cascade filter: flushing

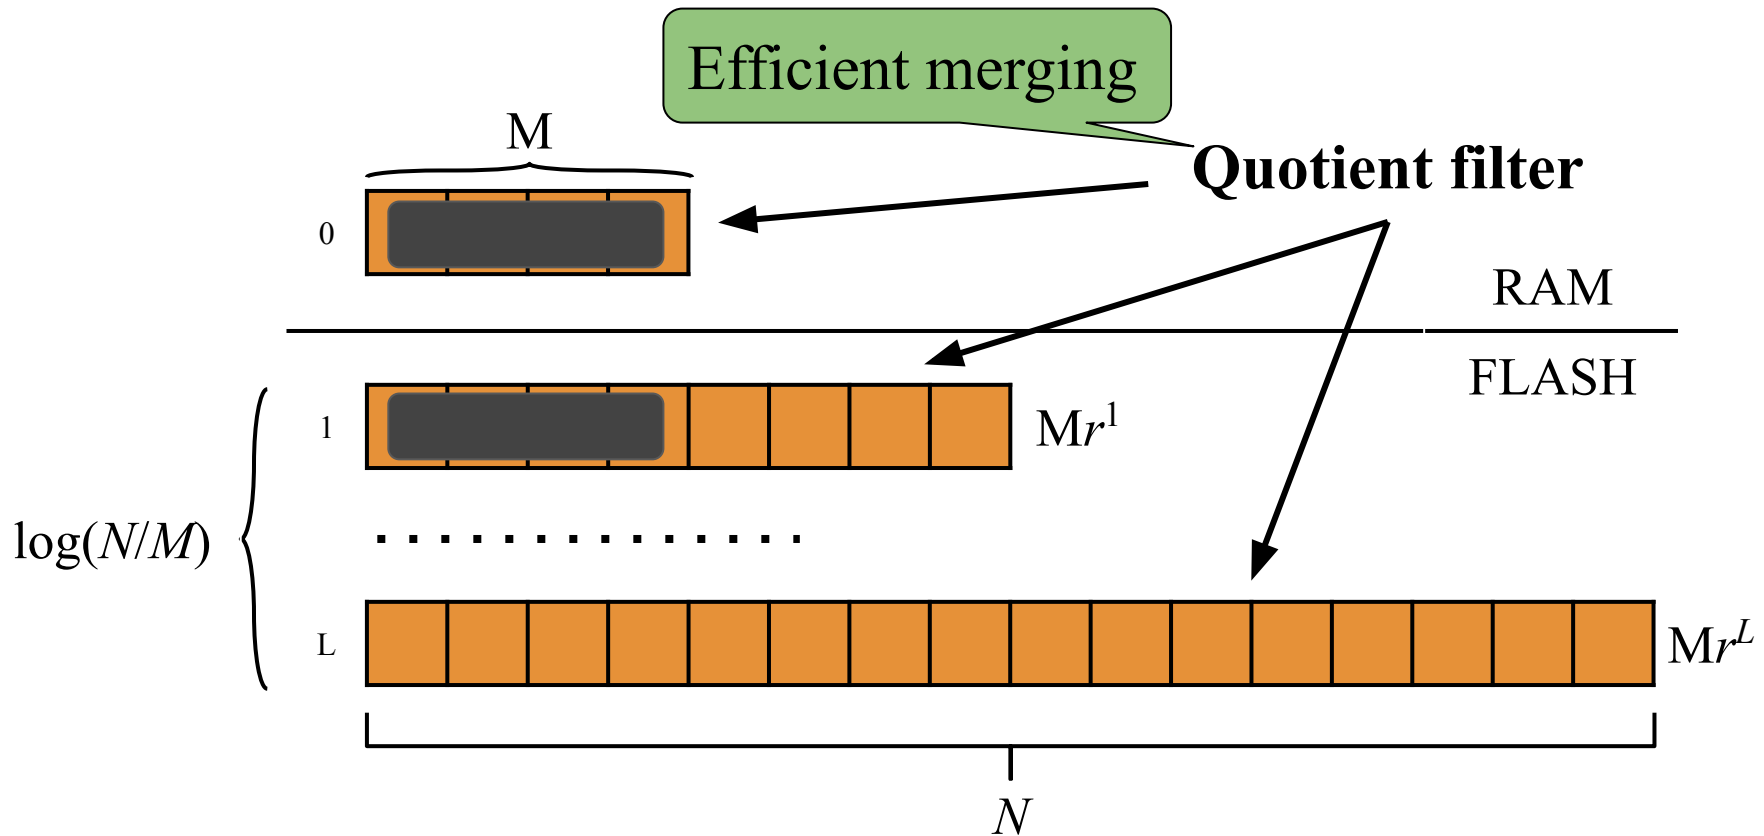
[Bender et al. '12, Pandey et al. '17]



When RAM is full, items are flushed to the smallest level on disk ***i*** with space to insert items in level **0** to ***i-1***

Cascade filter: flushing

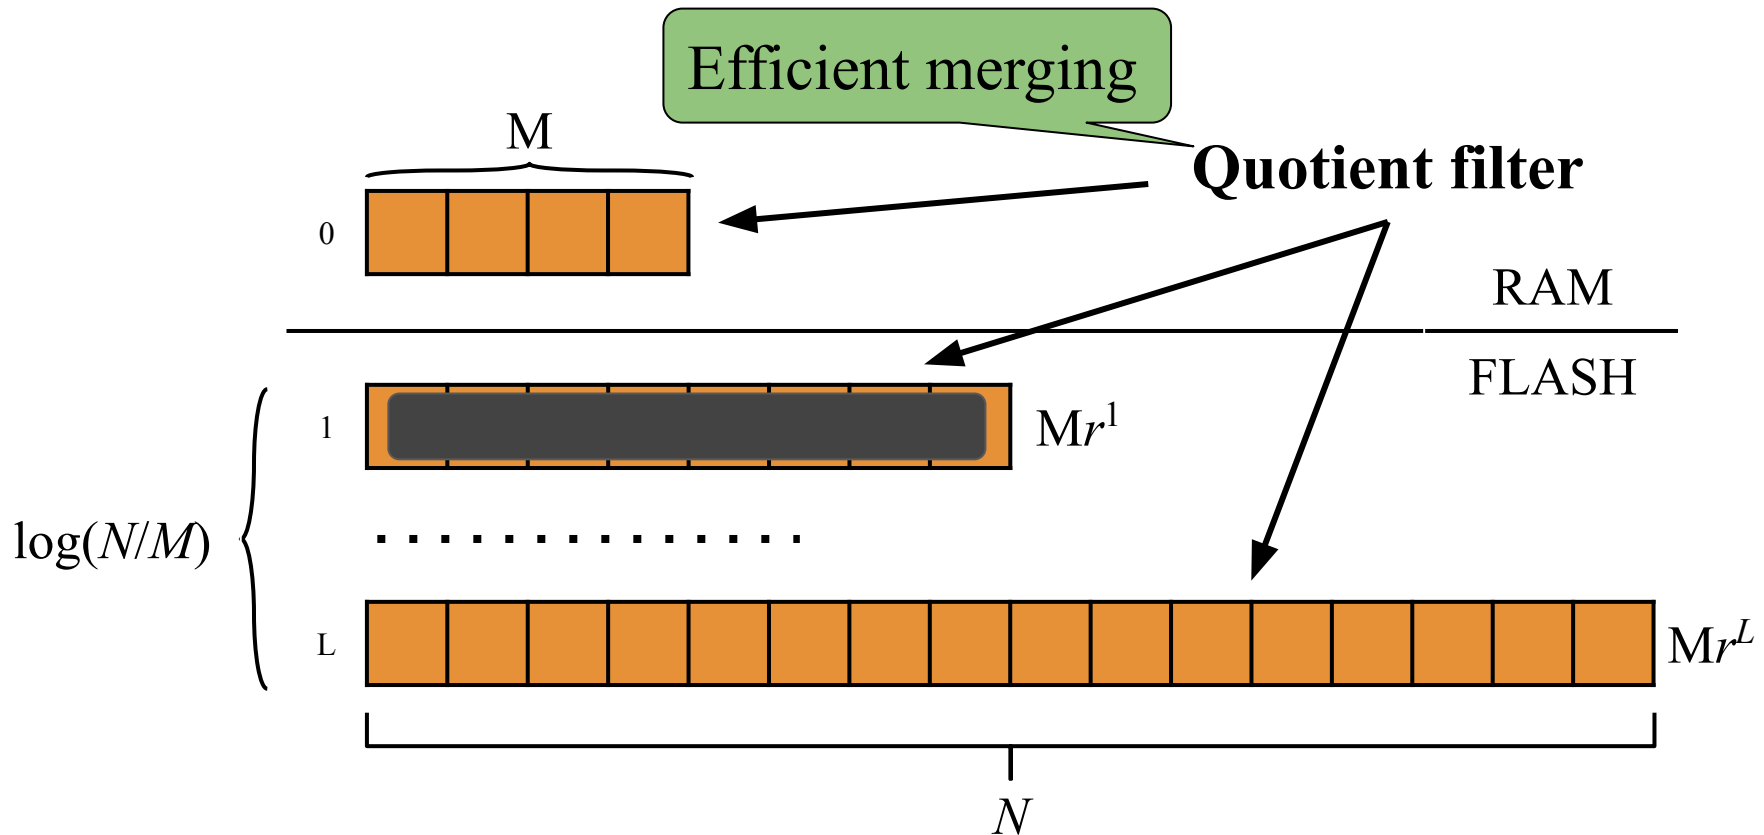
[Bender et al. '12, Pandey et al. '17]



When RAM is full, items are flushed to the smallest level on disk i with space to insert items in level **0** to $i-1$

Cascade filter: flushing

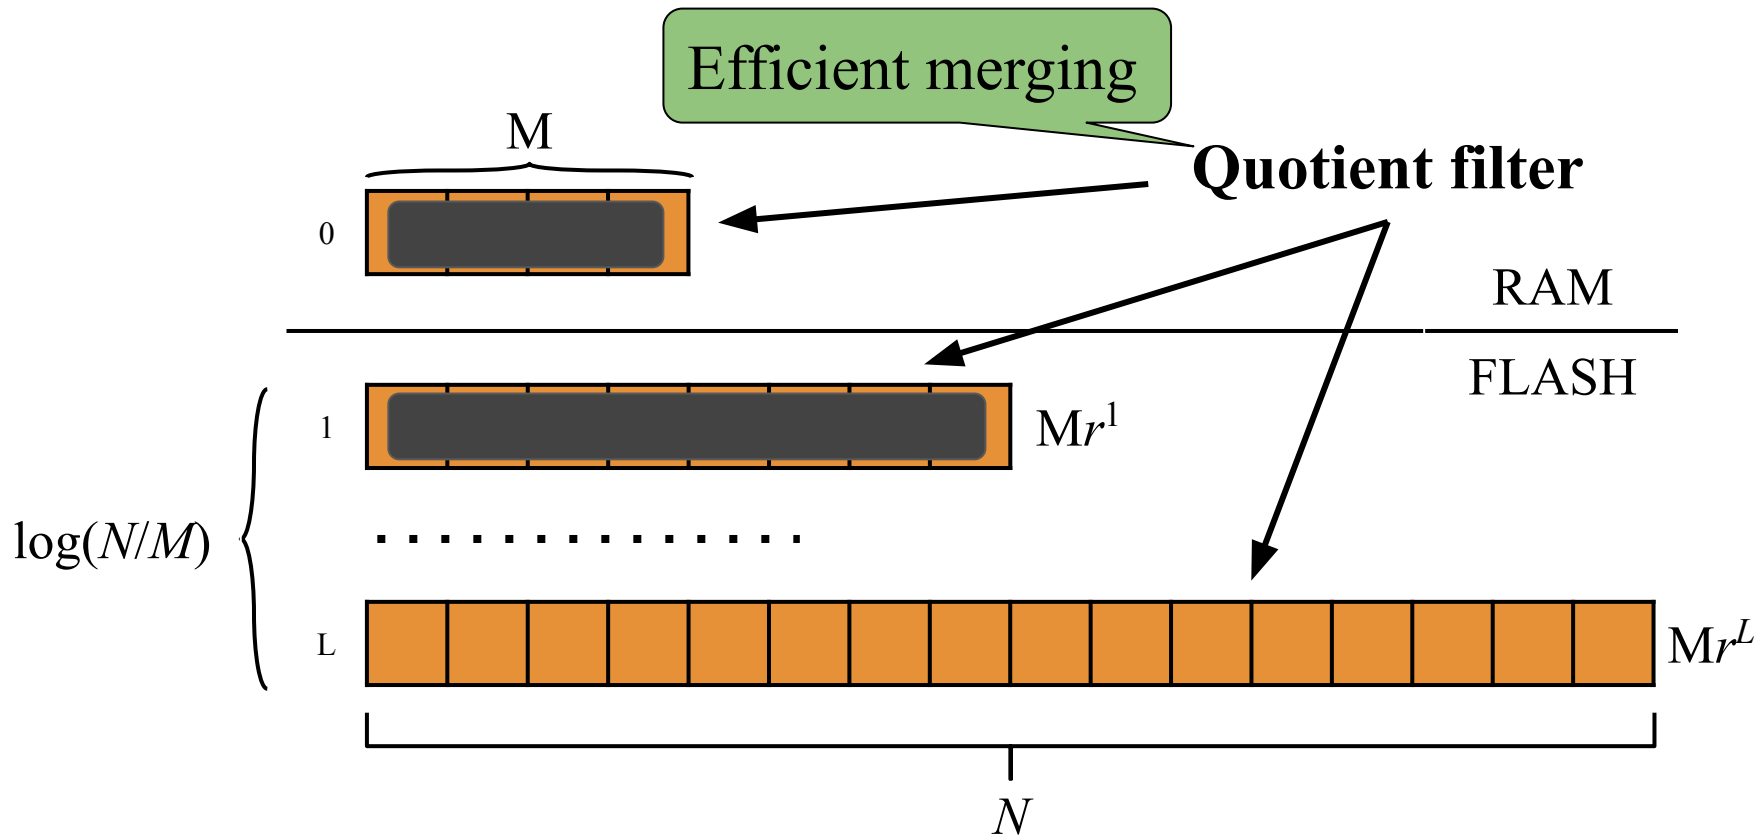
[Bender et al. '12, Pandey et al. '17]



When RAM is full, items are flushed to the smallest level on disk i with space to insert items in level **0** to $i-1$

Cascade filter: flushing

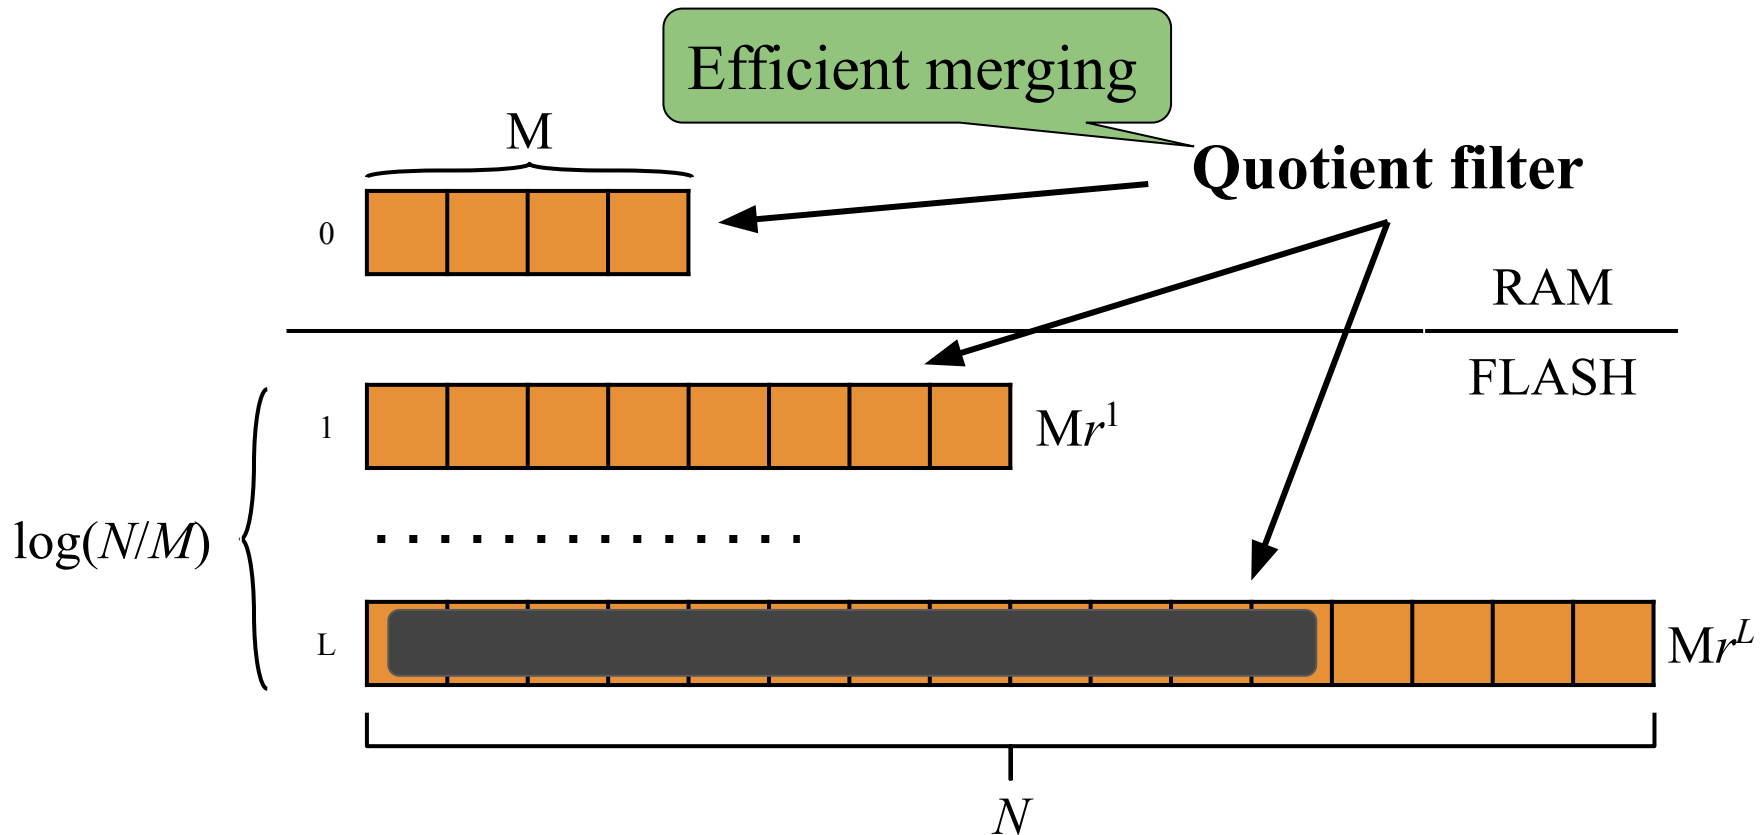
[Bender et al. '12, Pandey et al. '17]



When RAM is full, items are flushed to the smallest level on disk i with space to insert items in level **0** to $i-1$

Cascade filter: flushing

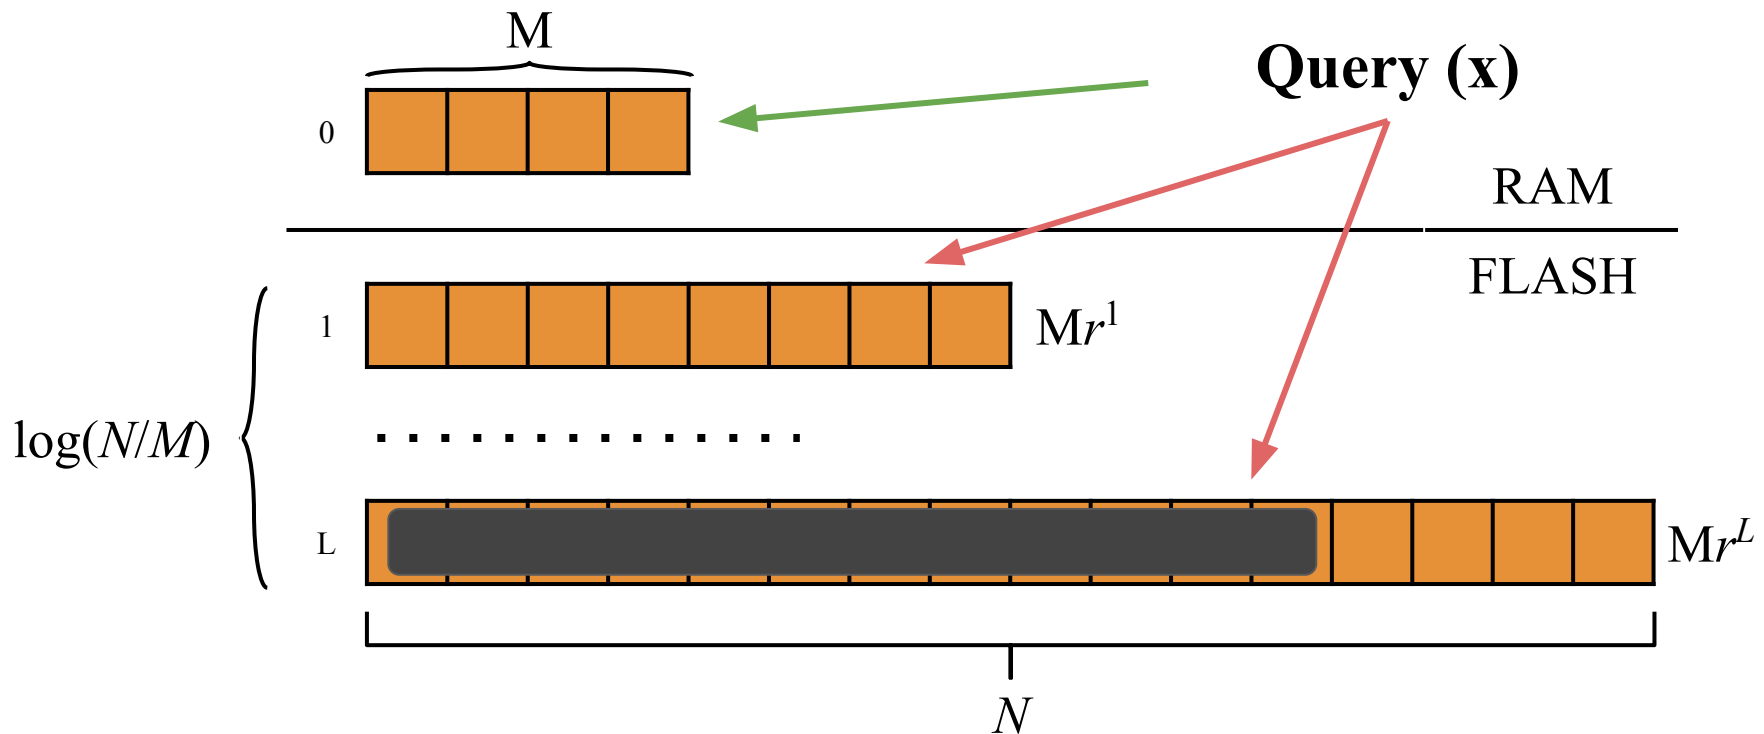
[Bender et al. '12, Pandey et al. '17]



When RAM is full, items are flushed to the smallest level on disk ***i*** with space to insert items in level **0** to ***i-1***

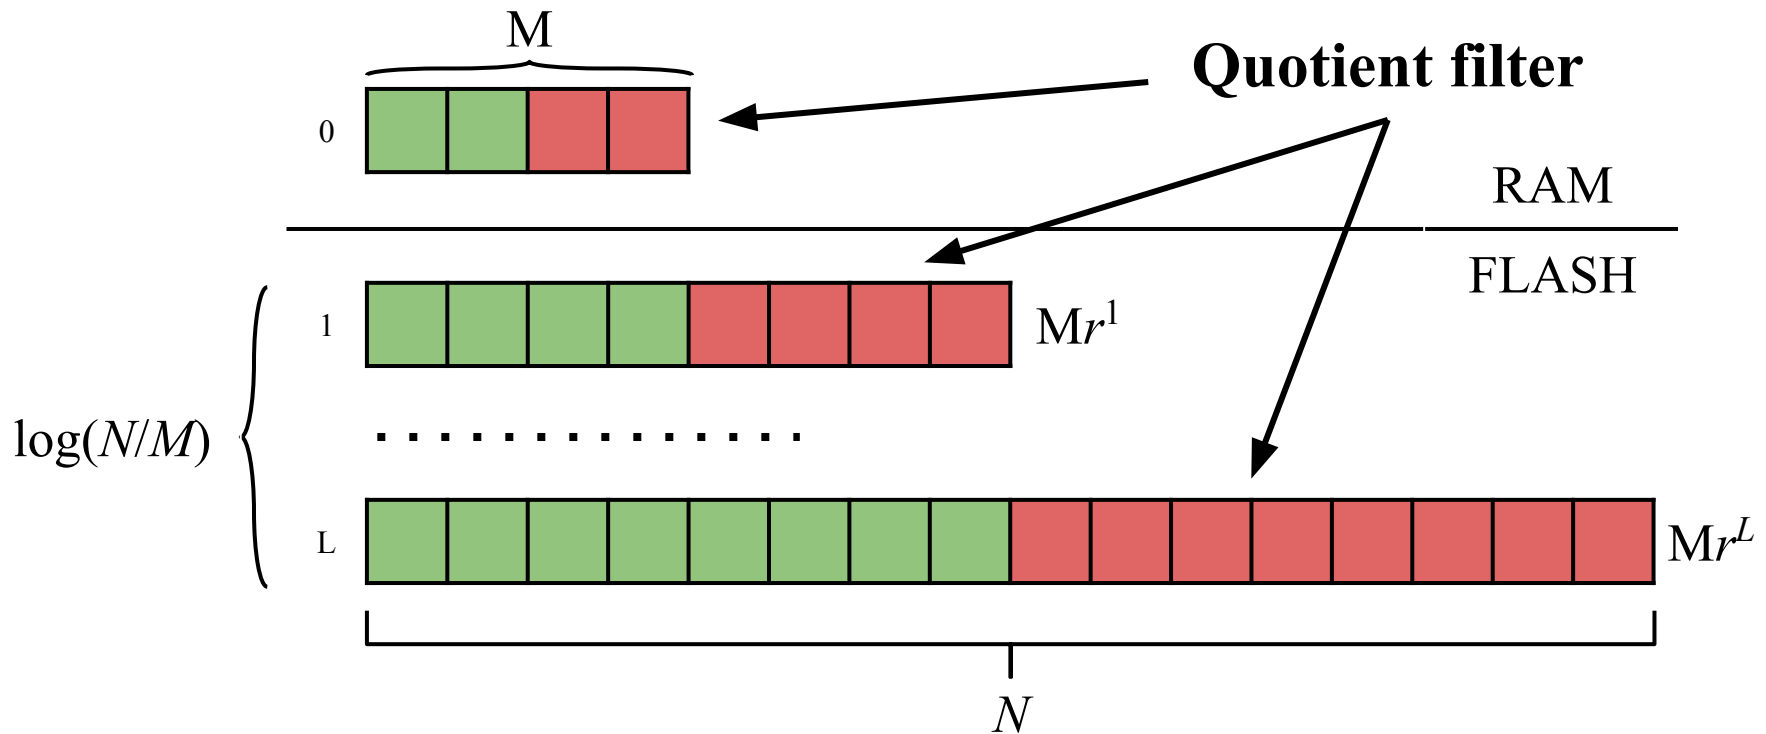
Cascade filter: query

[Bender et al. '12, Pandey et al. '17]



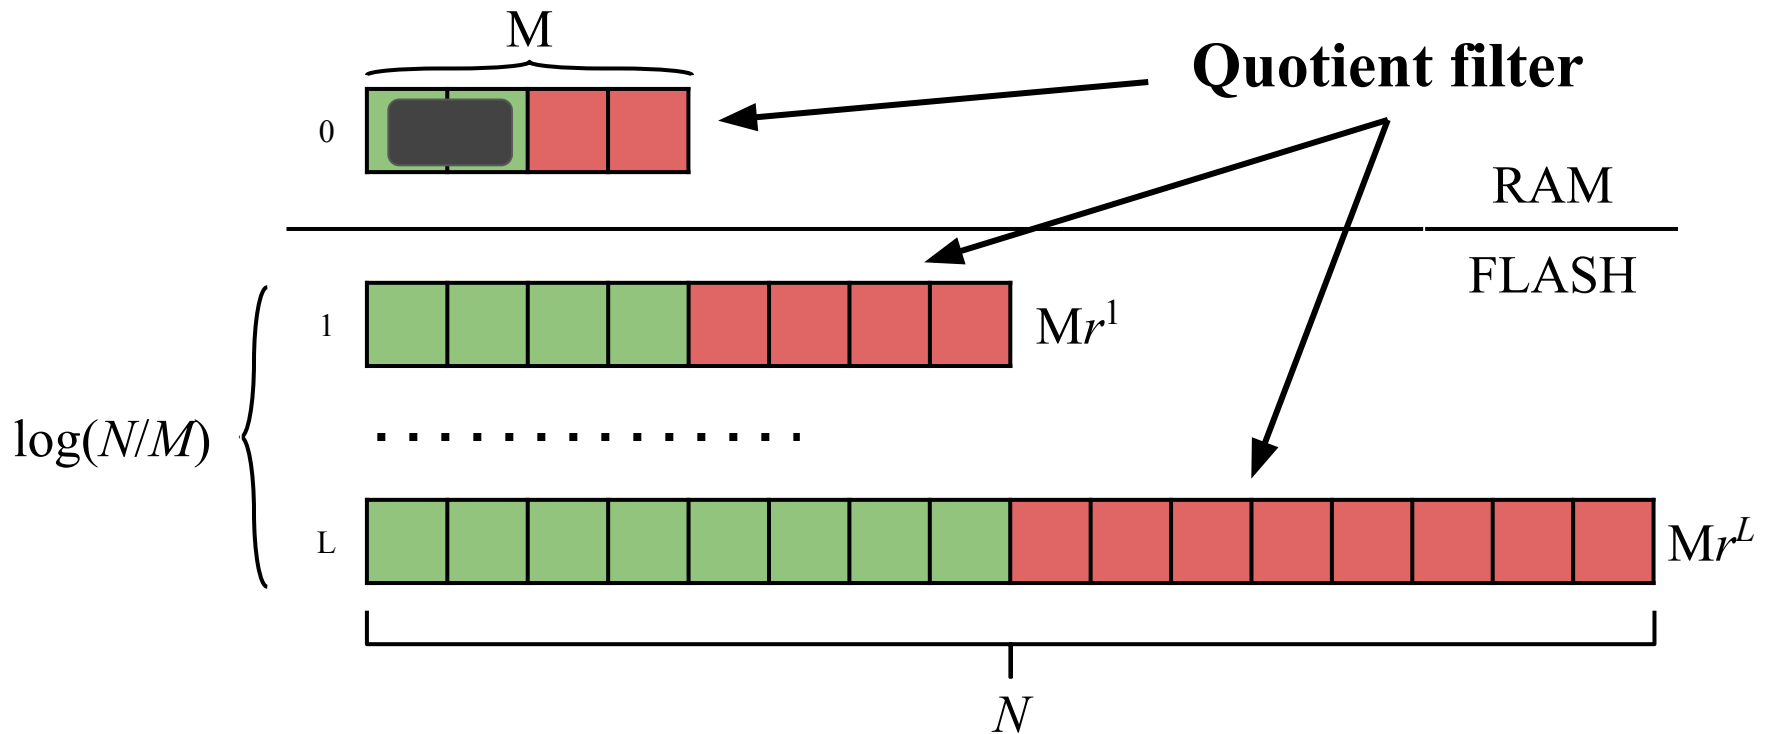
A query operation requires a lookup in each non-empty level

Time-stretch LERT




Divide each level into $1 + 1/\alpha$, equal-sized bins.

Time-stretch LERT



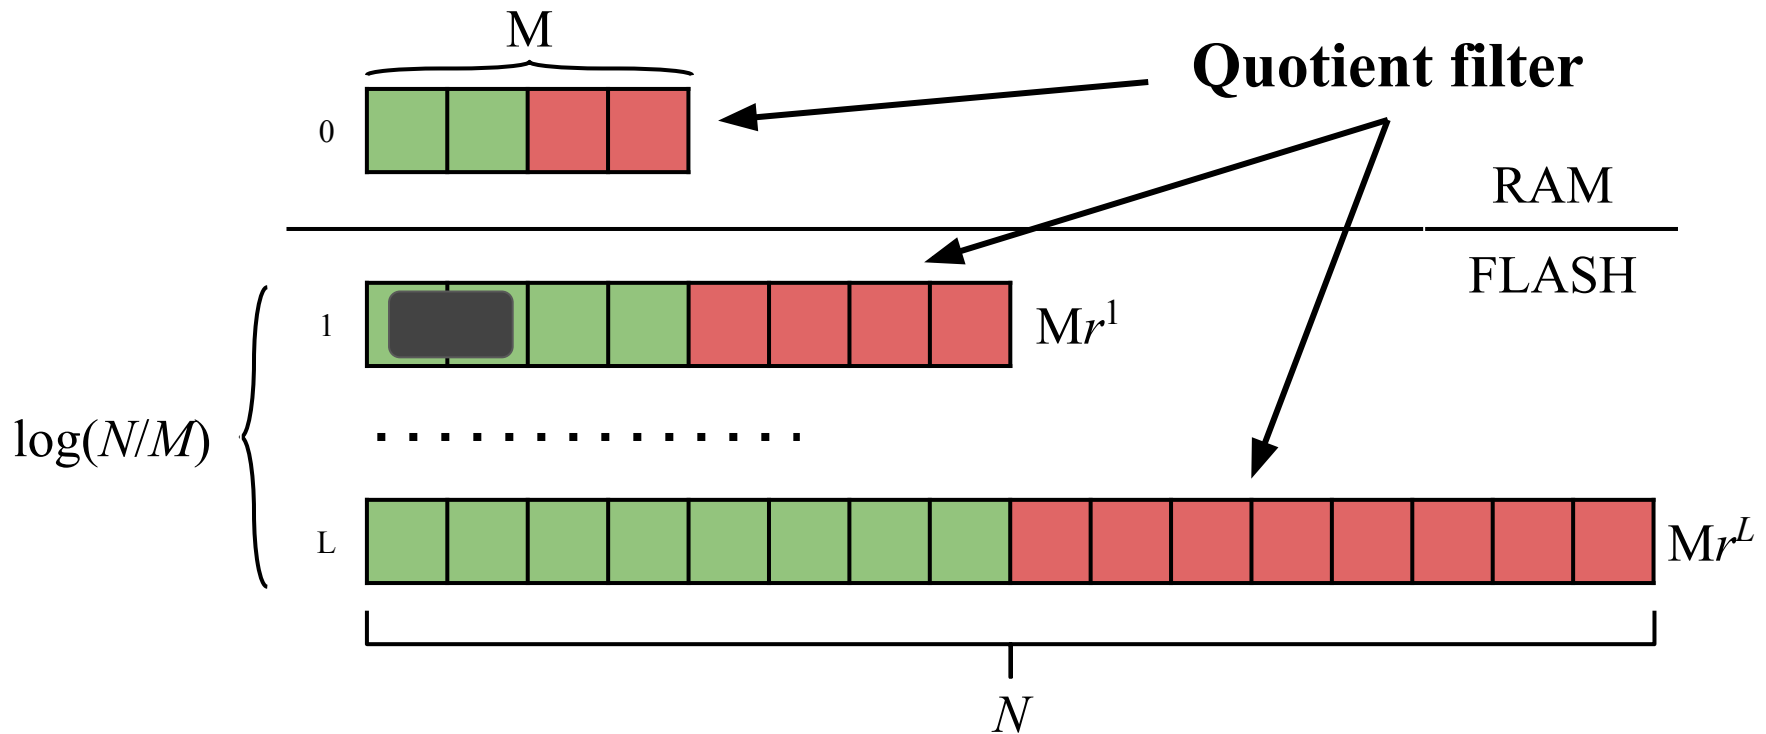
When a bin is full, items move to the adjacent bin

Time-stretch LERT



When a bin is full, items move to the adjacent bin

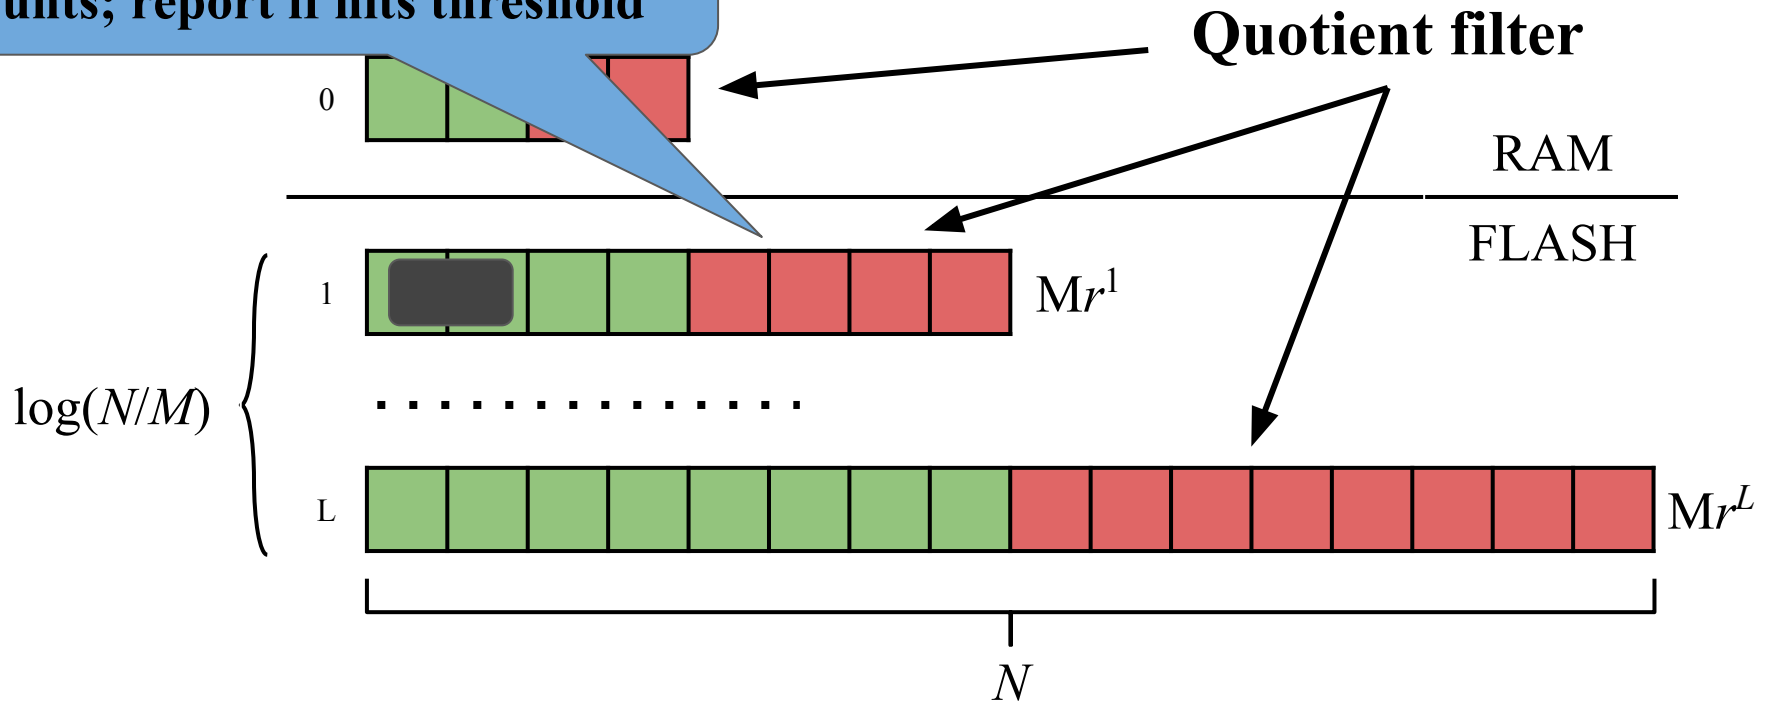
Time-stretch LERT



Last bin **flushed** to first bin of the next level

Time-stretch LERT

While flushing consolidate counts; report if hits threshold

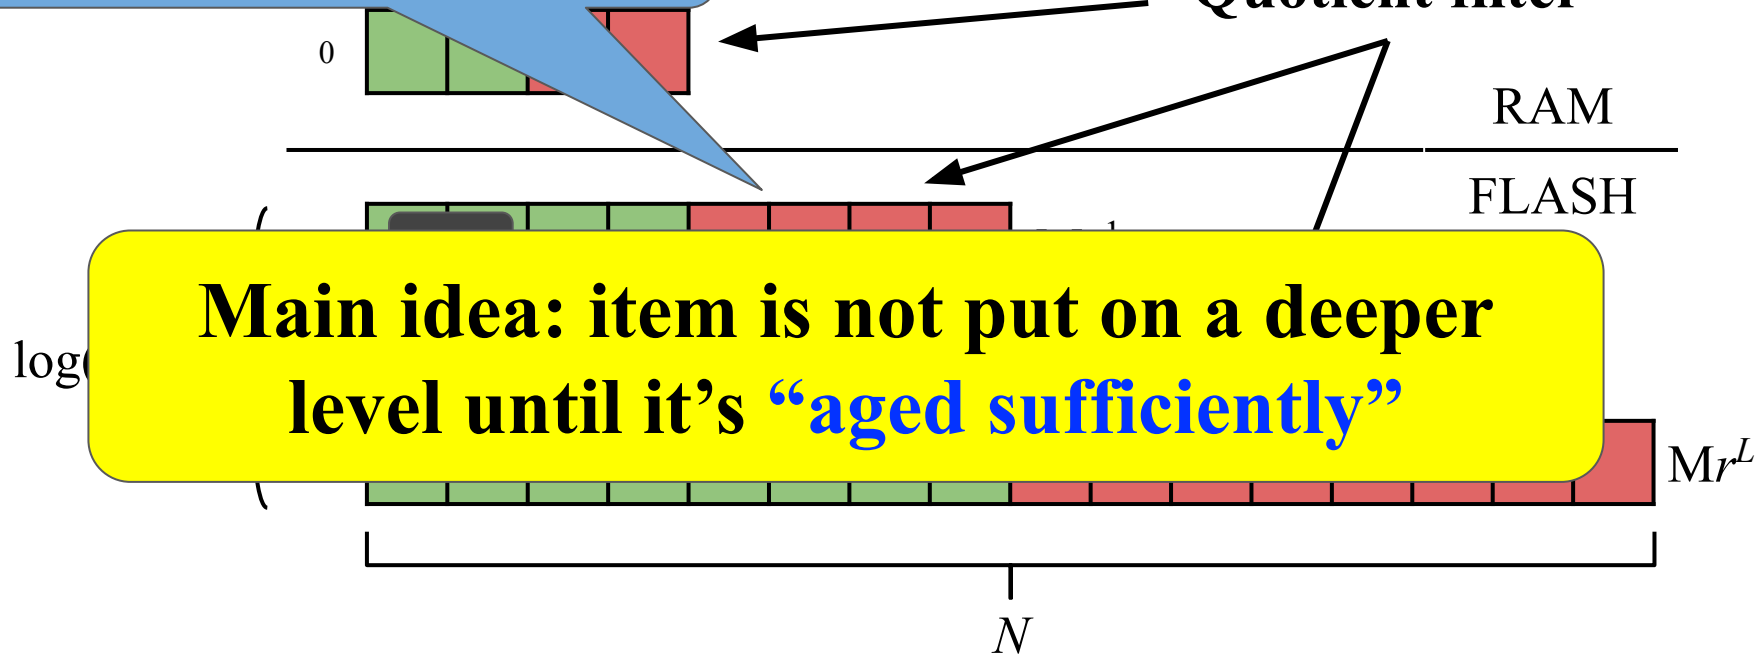


Last bin **flushed** to first bin of the next level

Time-stretch LERT

While flushing consolidate counts; report if hits threshold

Quotient filter



Last bin **flushed** to first bin of the next level

Time-stretch LERT I/O complexity

$$O \left(\left(\frac{\alpha+1}{\alpha} \right) \frac{1}{B} \log \frac{N}{M} \right)$$

Optimal insert cost for
Write-optimized data
structure

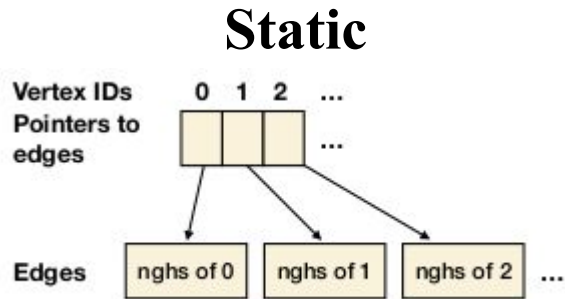
Time-stretch LERT I/O complexity

$$O \left(\left(\frac{\alpha+1}{\alpha} \right) \frac{1}{B} \log \frac{N}{M} \right)$$

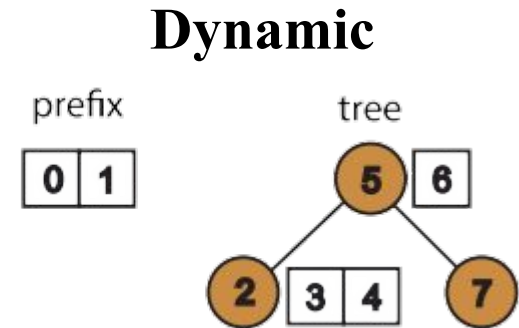
Extra cost because we only move one bin during a flush. Constant loss for constant α

Optimal insert cost for Write-optimized data structure

Trade-off 2: “One-size-fits-all” approach leaves performance on table



LIGRA [Shun & Blelloch '13]



ASPEN [Dhulipala et al. '19]

	LIGRA	ASPEN
add_edge	$O((E + V)/B)$	$O(\log V + c^2 \log(\deg(u))/B)$
get_neighbors	$O(\deg(u)/B)$	$O(\log V + \deg(u)/B + \deg(u)/c)$

Neighbor access requires at least *two cache misses*

For dynamic, all operations have a *log factor*

Trade-off 2: “One-size-fits-all” approach leaves performance on table



Static → Fast computations; no updates
Dynamic → Slower computations; updates


	LIGRA	ASPEN
add_edge	$O((E + V)/B)$	$O(\log V + c^2 \log(\deg(u))/B)$
get_neighbors	$O(\deg(u)/B)$	$O(\log V + \deg(u)/B + \deg(u)/c)$

Neighbor access requires at least *two cache misses*

For dynamic, all operations have a *log factor*

Real world graphs are often skewed

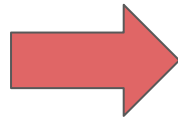
High variance in the degree distribution



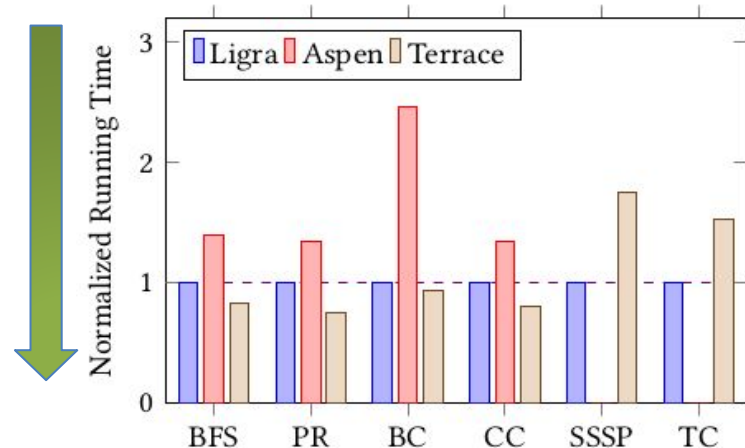
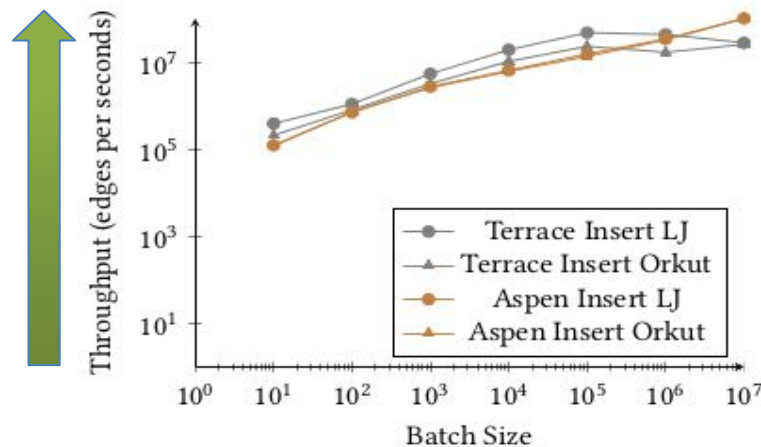
- Dynamic partitioning of vertices based on the degree
- Separate structures for each partition to minimize cache misses

Dynamic partitioning + hierarchical structure

High variance in the degree distribution

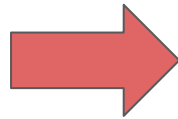


- Dynamic partitioning of vertices based on the degree
- Separate structures for each partition to minimize cache misses



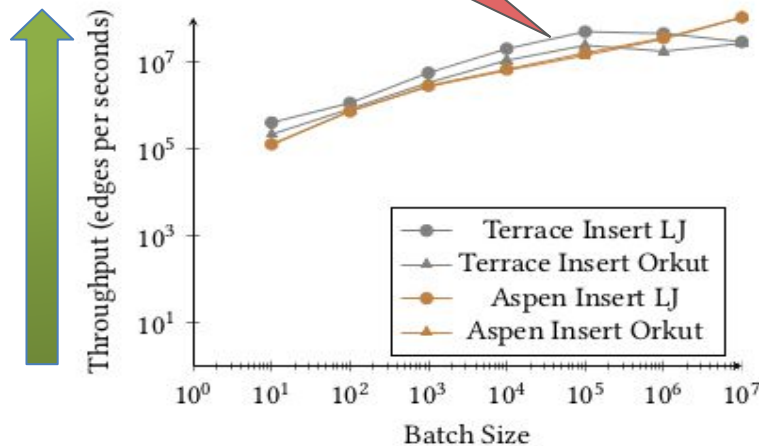
Dynamic partitioning + hierarchical structure

High variance in the degree distribution

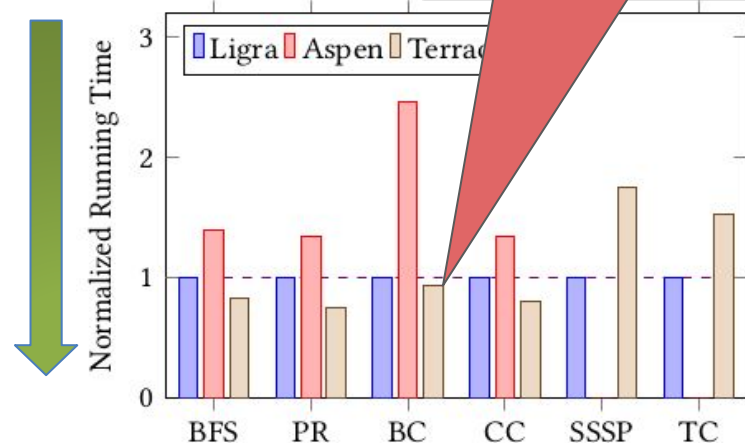


- Dynamic partitioning of vertices based on the degree
- Separate structures for each partition to minimize cache misses

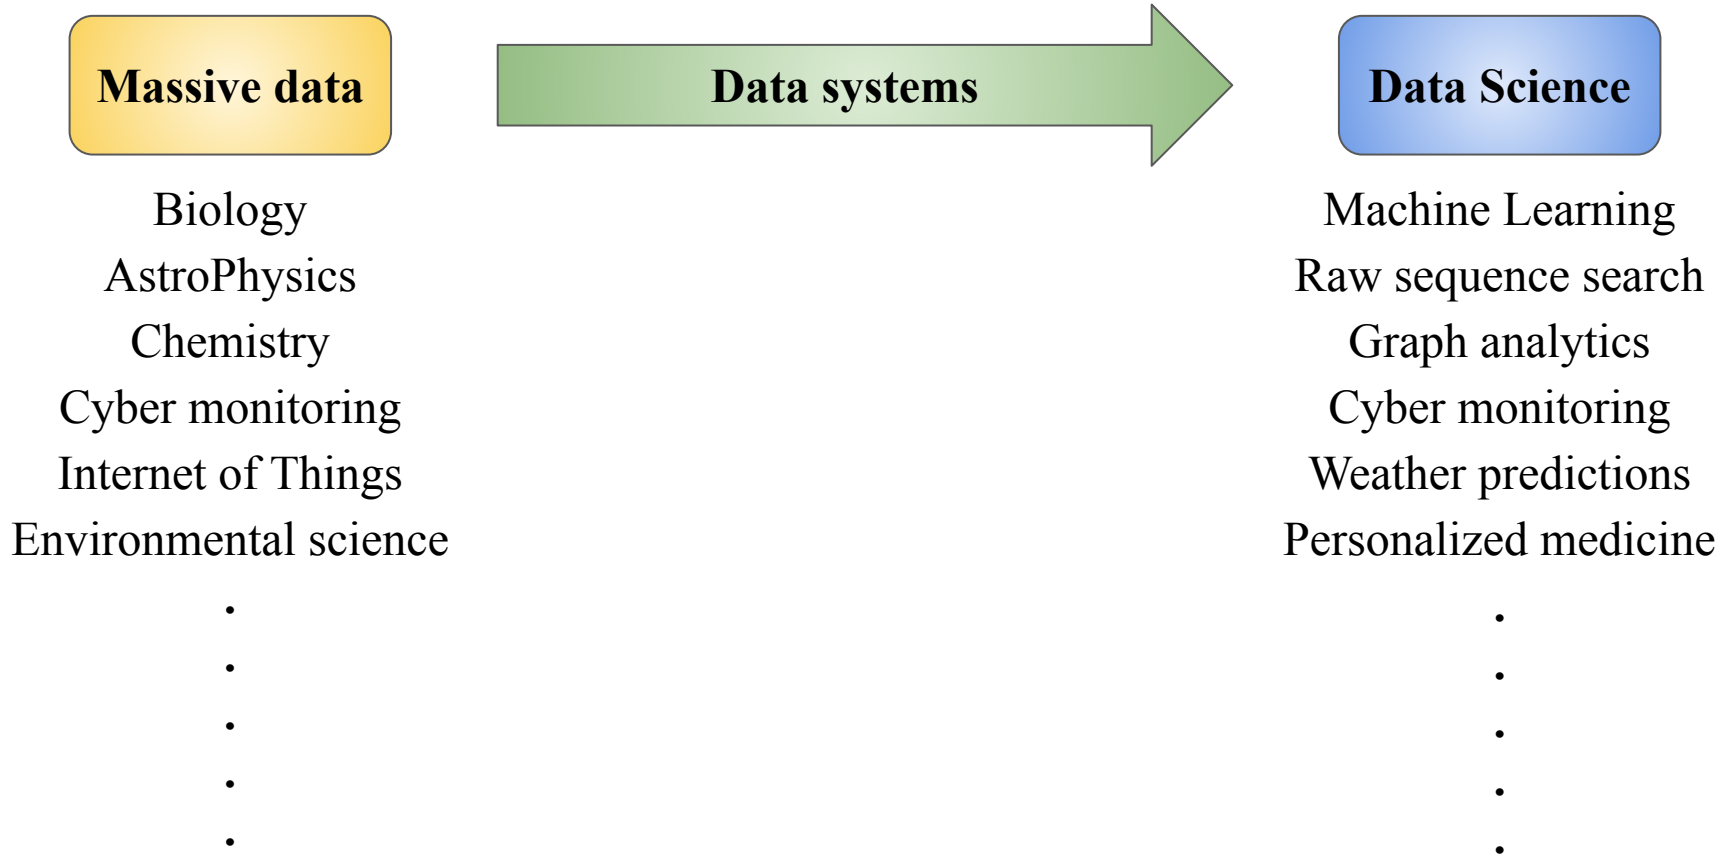
**Terrace:
Fast updates**



**Terrace:
Faster computations**



Scalable data systems → Scalable data science



My goal as a researcher is to build *scalable data systems* to *accelerate* and *scale data science* applications

Our contribution



**Combine streaming and EM algorithms to solve
real-time event detection problem**